

Contents

Summary	3
Contributor Bios	4
Acknowledgements	5
Introduction (Vivek Srikrishnan and Klaus Keller)	6
Chapter #0: Introduction to Python (Giacomo Marangoni and Vivek Srikrishnan)	7
Chapter #1: Values and Mental Models in Decision Support (Casey Helgeson)	8
Chapter #2: Bayesian Inference and Markov chain Monte Carlo Basics (Kelsey L. Ruckert, Tony E. Wong, Benjamin Seiyon Lee, Yawen Guan, and Murali Haran)	9
Learning objectives	9
Introduction	9
Why does MCMC work?	10
Detailed balance equation	11
Metropolis-Hastings Algorithm	11
How does MCMC work?	12
Tutorial	14
Exercise	20
Questions	21
References	21
Chapter #3: A Calibration Problem and Markov chain Monte Carlo (Kelsey L. Ruckert, Tony E. Wong, Yawen Guan, Murali Haran, and Patrick J. Applegate)	22
Learning objectives	22
Introduction	22
Tutorial	22
Setting up the prior information and the likelihood function	23
Bayesian Inference using MCMC	26
Checking the acceptance rate	26
Testing for convergence	27
Analyzing MCMC output	31
Exercise	31
Questions	33
Appendix	33
References	34
Chapter #4: Applying Markov chain Monte Carlo to sea-level data (Kelsey L. Ruckert, Tony E. Wong, Yawen Guan, and Murali Haran)	35
Learning objectives	35
Introduction	35
How do we account for the error structure?	35
Tutorial	37
Characterizing the error structure	38
Setting up the prior information	40
Running MCMC	41
Determine the burn-in period	41
Hindcasting and projecting sea-level rise	43
Exercise	43
Questions	44

Appendix	45
References	46
Chapter #5: Global sensitivity analyses (Tony E. Wong, Vivek Srikrishnan and Klaus Keller)	48
Chapter #6: Decision-making under uncertainty (Vivek Srikrishnan and Klaus Keller)	49
Chapter #7: Climate policy and the DICE model (Patrick J. Applegate, Gregory G. Garner, Richard B. Alley, and Klaus Keller)	50
Learning objectives	50
Introduction	50
An important note	51
Tutorial	51
Running DICE in optimization mode and examining its time-varying output	52
Accounting for uncertainty in the climate sensitivity	54
Exercise	55
Questions	55
References	56
Chapter #8: Multi-objective decision analysis with a lake management problem (Caitlin Spence, Vivek Srikrishnan, and Giacomo Marangoni)	57
Chapter #9: Introduction to multi-objective evolutionary algorithms and multi-objective optimization (Vivek Srikrishnan, Caitlin Spence, Giacomo Marangoni, and Patrick Applegate)	58
Chapter #10: Vulnerability analysis and robust decision-making (Vivek Srikrishnan, Caitlin Spence, and Giacomo Marangoni)	59
Chapter #11: Principles of data visualization (Vivek Srikrishnan, Kelsey L. Ruckert, and Klaus Keller)	60
Chapter #12: Visualizing uncertainty using sea-level rise and storm surge (Vivek Srikrishnan, Kelsey L. Ruckert, and Klaus Keller)	61
GNU GENERAL PUBLIC LICENSE Version 3 (Free Software Foundation, Inc.)	62

Summary

This e-textbook presents a series of chapters with code examples and exercises in R and Python that introduce advanced aspects of risk analysis and decision support. These topics are intended for graduate students and practitioners.

Contributor Bios

Vivek Srikrishnan is an Assistant Research Professor at Penn State. He is interested in how communities and regions manage climate risks. This research focuses on decision-making about climate change adaptation and mitigation under coupled social, economic, and Earth system uncertainties. Vivek received his Ph.D and master's from Penn State in energy and mineral engineering.

Klaus Keller is a Professor of Geosciences at Penn State, where he directs the Center for Climate Risk Management as well as the research network for Sustainable Climate Risk Management (<http://scrimhub.org>). Before joining Penn State, he worked as a research scientist and lecturer at Princeton University and as an engineer in Germany. Professor Keller graduated from Princeton with a Ph.D. in civil and environmental engineering. He received master's degrees from M.I.T. and Princeton as well as an engineer's degree from the Technische Universität Berlin. His research addresses two interrelated questions. First, how can we mechanistically understand past and potentially predict future changes in the climate system? Second, how can we use this information to design sustainable, scientifically sound, technologically feasible, economically efficient, and ethically defensible climate risk management strategies? He analyzes these questions by mission-oriented basic research covering a wide range of disciplines such as Earth system science, economics, engineering, philosophy, decision science, and statistics. Klaus' e-mail address is klaus@psu.edu, and his Web site is at <http://www3.geosc.psu.edu/~kzk10/>.

Acknowledgements

This work was supported by the National Science Foundation through the Network for Sustainable Climate Risk Management (SCRiM) under NSF cooperative agreement GEO-1240507. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation. Other support was provided by the Center for Climate Risk Management and the Rock Ethics Institute.

Katerina Kostadinova designed the cover using a photograph provided by Frank Erickson.

Randy Miller created scripts for merging the parts of the book into a single `.pdf` file and created the Github repository.

Vivek Srikrishnan thanks his colleagues and the participants in the 2016 SCRiM Summer School for helping to shape the scope of this book.

Klaus Keller thanks all the students from the Risk Analysis course, the members of the Keller research group, and his SCRiM collaborators for invaluable inputs. Special thanks go to his family, many mentors, and friends for support.

Introduction (Vivek Srikrishnan and Klaus Keller)

Chapter #0: Introduction to Python (Giacomo Marangoni and Vivek Srikrishnan)

Chapter #1: Values and Mental Models in Decision Support (Casey Helgeson)

Chapter #2: Bayesian Inference and Markov chain Monte Carlo Basics (Kelsey L. Ruckert, Tony E. Wong, Benjamin Seiyon Lee, Yawen Guan, and Murali Haran)

Learning objectives

After completing this exercise, you should be able to

- describe what the detailed balance equation is
- describe what are the three conditions
- explain why the Metropolis-Hastings algorithm fulfills the conditions and works

Introduction

In earlier chapters, we examined how a statistical technique called the bootstrap can be used to determine how confident we are about our estimates of uncertain values. In particular, we looked at how to determine whether a coin is fair and how sure we can be about our estimates of future sea-level change.

The bootstrap is an example of what is called a *frequentist* statistical technique. Frequentist techniques assume that the probability of an event is the proportion of the times the event will occur (if it was an infinite number of trials and a probability $0 < p < 1$, then we would see an infinite number of “successful” outcomes). For example, we know that the probability that a fair coin will come up heads on any individual flip is 0.5. In a frequentist interpretation, that probability implies that we can observe a very large number of coin flips, in which we count the number of times the coin comes up heads. That number of times, divided by the number of flips, is the *frequency* with which the coin comes up heads.

Of course, in many situations, we cannot perform lots of random trials to determine the probabilities of different outcomes. The sea-level rise problem is one example. We want to know how much sea level will rise in the future. Because the data and physical models are imperfect, we cannot be sure of the exact answer. In a frequentist framework, we would try to ascertain how sea-level rise would vary under multiple hypothetical replications of the state of the world (multiple alternative worlds). This is how we would determine the probabilities of different outcomes.

Bayesian techniques provide us with an alternative way of viewing this problem. In a Bayesian framework, we have a preexisting estimate of the probability of different outcomes that is based on our past experiences and our beliefs about the situation in question. We then make observations and update the probability estimates based on those observations. This procedure, which is called *Bayesian updating*, is perhaps similar to how people often make decisions. New information leads to changed opinions. The entire process of defining the prior probability distribution of different outcomes or physical model parameters and then using Bayesian updating to update our beliefs about the outcomes is called *Bayesian model calibration*.

It is important to point out that the use of the word “model” can refer to two different things, that is, the statistical model (which is Bayesian in our case) and the physical model of the system of interest (e.g., a global sea-level model). For simplicity and ease of understanding, we will use “statistical” or “physical” when referring to both kinds of models.

In Bayesian model calibration, physical model parameters are considered to be random variables. Our knowledge of the parameters (before any data are observed) is represented by a *prior* probability distribution. Observations may be used to inform estimates of which parameter values are more or less likely. The probability model for observations provides a distribution on observations for a particular parameter setting, that is, as we vary the value of the parameters, the probability distribution changes. To fix ideas, think of the mean and variance parameters of a normal distribution — as we vary the values of the mean and the variance, the normal distribution for the observation changes. The probability model therefore provides a probability distribution for random variables for a particular parameter value. A *likelihood function* helps solve the

inverse problem — it is useful for providing information about the parameters *given the observations*. It is obtained from the probability distribution by plugging in the observations into the function. The likelihood function is therefore a function of just the parameters — this includes both the statistical and physical model parameters. Direct sampling from the posterior distribution is typically impossible because the posterior is only known up to a constant, and in many cases this distribution is intractable. In order to obtain samples from the posterior, one approach is to employ a *Markov chain Monte Carlo* (MCMC) sampling technique.

MCMC is an algorithm used to simulate random variables or “draw samples” from a given probability distribution by constructing a *Markov chain* based on the distribution. For Bayesian inference, the distribution of interest is the posterior distribution. A Markov chain is a sequence of random variables where each successive value in the sequence depends on the current value of the sequence. The *Metropolis-Hastings* algorithm is used to construct the Markov chain so its “stationary distribution” is the distribution of interest. What this means is that the Markov chain satisfies a theoretical property which allows us to use sample means of the Markov chain to approximate the mean of the posterior distribution (i.e., using the least squares error in the likelihood function). For instance, if we want to approximate the mean of the posterior distribution, we just need to take an average of the Markov chain samples. If we want to approximate the correlation between two random variables in a joint distribution, we can take the sample correlation between the two samples in the Markov chain.

MCMC’s popularity and prominence are due to its generality as it may be used to draw samples from high-dimensional, complicated probability distributions for which sampling algorithms may not generally exist. MCMC is useful where the goal is to use posterior distributions of parameters in a physical model to summarize the information about the physical model parameters that are contained in a given data set. Because posterior distributions tend to be complicated, MCMC is often one of the few generally applicable approaches that may be used to approximate various characteristics of the posterior distribution. Once samples from the posterior distribution are generated, one can calculate “best” (“point”) estimates for the parameters by using attributes like posterior means and also represent uncertainties for the parameters of interest by using “credible intervals” (Bayesian analogues to frequentist confidence intervals) based on the samples. Credible intervals can be estimated with “quantiles”. In the next chapter, this process of estimating the credible intervals will be described. Bayesian methods are useful in climate science because they can easily integrate multiple sources of information (e.g., physical models, multiple data sets, and complex sources of error) into a single framework. MCMC then provides a standard algorithm for approximating the resulting posterior distribution, thereby conveniently incorporating multiple sources of uncertainty, say from various data and assumptions, when providing conclusions about the model (physical and statistical) parameters.

MCMC is a complex concept, so we broke up the section on MCMC into three chapters. Each chapter will build off of the previous one. In this chapter, the goal is to describe how MCMC works and why it works. To do this you will code your own MCMC with the Metropolis-Hastings algorithm. The overall goal of these three chapters are to:

1. understand the basics of MCMC
2. understand the output and what MCMC convergence means
3. how to apply and problem solve applying MCMC to a “real” model and “real” data

Why does MCMC work?

Here we briefly describe some of the basic concepts of MCMC and why it works. For more details, we recommend consulting chapter 6 from Wood (2015), the YouTube video called, “(ML 18.6) Detailed balance (a.k.a. Reversibility)”, Detailed balance (2017), and Gilks (1997). This section was adapted from those references.

Detailed balance equation

To understand how and why MCMC works, one must learn about *detailed balance equations*. Detailed balance or reversibility is when a process—say the probability density distribution π on some set of states Y —satisfies the detailed balance equations with respect to a *transition probability* distribution T . The transition probability is the probability of changing from one parameter value to another value in a single move. If π satisfies detailed balance with respect to T , then this implies π has a stationary distribution such that

$$\pi(x)T(x, y) = \pi(y)T(y, x),$$

where $T(x, y)$ is the transition probability of moving from state x (current value) to state y (proposed value), and $\pi(x)$ and $\pi(y)$ are the equilibrium probabilities of being in states x and y , respectively. In other words, for all y ,

$$\pi(y) = \sum_x \pi(x)T(x, y) = \sum_x \pi(y)T(y, x) = \pi(y) \sum_x T(y, x) = \pi(y).$$

Additionally, detailed balance also implies that the process must be reversible. In other words, the Markov chain looks the same moving forward as it does moving backwards. For instance, a Markov chain, MC, satisfies detailed balance when, $(Y_0, \dots, Y_n) \sim MC(\pi, T) \Rightarrow (Y_n, \dots, Y_0) \sim MC(\pi, T)$. Therefore, detailed balance implies that there is no net flow of probability such that

$$T(x, y)T(y, z)T(z, x) = T(x, z)T(z, y)T(y, x).$$

Metropolis-Hastings Algorithm

Next, we show how the Metropolis-Hastings algorithm satisfies the detailed balance equation. Given that we have two values of the Markov chain (x and y), $k(x, y)$ is the *transition kernel* of moving from x to y , and $k(y, x)$ is the transition kernel of moving from y to x , we want to show that

$$(1) \quad \pi(x)k(x, y) = \pi(y)k(y, x).$$

In the Metropolis-Hastings algorithm, the transition kernel, $k(y, x)$ of moving from y to x is defined as

$$(2) \quad k(y, x) = \alpha(y, x)q(y, x),$$

where $q(x, y)$ is the proposal distribution and $\alpha(y, x)$ is the acceptance probability for the proposal. The acceptance probability is defined as

$$(3) \quad \alpha(y, x) = \min \left\{ 1, \frac{\pi(x)q(x, y)}{\pi(y)q(y, x)} \right\}.$$

In equation (1), if $x = y$, then the detailed balance equation is satisfied. If $x \neq y$, then the left hand side of the equation (LHS) gives us

$$(4) \quad \pi(x)k(x, y) = \pi(x)q(x, y)\alpha(x, y).$$

Without loss of generality, we assume that $\pi(y)q(y, x) > \pi(x)q(x, y)$. Then, by equation (3), the LHS becomes:

$$(5) \quad \pi(x)k(x, y) = \pi(x)q(x, y)\alpha(x, y) = \pi(x)q(x, y) \times 1.$$

The right hand side (RHS) of equation (1) becomes

$$\pi(y)k(y, x) = \pi(y)q(y, x)\alpha(y, x) = \pi(y)q(y, x) \times \frac{\pi(x)q(x, y)}{\pi(y)q(y, x)} = \pi(x)q(x, y) = \text{LHS}.$$

Hence, the Markov chain defined by the all-at-once Metropolis-Hastings algorithm satisfies the detailed balance equation with respect to the stationary distribution π .

How does MCMC work?

A Markov chain is determined by three key elements: the *parameter space*, the initial distribution, and the transition probability distribution. Parameter space is the set of all possible parameter value combinations that satisfy the specified constraints. For MCMC, where Monte Carlo simply means random samples, we generate a random parameter value from the initial distribution, then move to the next parameter value iteratively based on the transition probability. Under certain conditions, the draws will eventually reach an equilibrium probability distribution. That is, statistical properties of the distribution do not change as new random values are generated, thus showing evidence of *convergence*.

For example, suppose we are interested in the weather (state) on any day (this example is adapted from Example 11.1 in Grinstead and Snell (2006)). Suppose the weather can be either rainy, foggy, or sunny; thus, the parameter space contains three possible values, “rainy”, “foggy”, and “sunny”. Based on past experience, we know in this example that there are never two sunny days in a row and only half of the time a sunny day will occur after a foggy or rainy day. We also know there is an even chance of having two foggy days in a row and two rainy days in a row. This information can be represented as a *transition matrix*,

$$p = \begin{matrix} & \text{FOGGY} & \text{SUNNY} & \text{RAINY} \\ \text{FOGGY} & \left[\begin{array}{ccc} 1/2 & 1/4 & 1/4 \\ 1/2 & 0 & 1/2 \\ 1/4 & 1/4 & 1/2 \end{array} \right] \\ \text{SUNNY} & \\ \text{RAINY} & \end{matrix},$$

where each row represents the current state of the weather; a foggy day (first row), sunny day (second row), and rainy day (third row). The columns represent the next state of the weather; a foggy day (first column), sunny day (second column), and rainy day (third column). For example, if today is foggy, then the probability of the event that tomorrow is also foggy is $p_{1,1}$ or 0.5. Note that given the state of the weather today, the sum of the probabilities of the parameter values tomorrow must necessarily equal one.

We are going to formulate this example in R. The transition matrix can be set up by first creating vectors of the transition probabilities followed by using the matrix command. Typing `help(matrix)` in the Console window will describe the arguments of the function.

```
# Set up the known probabilities based on prior knowledge.
# Define the transition matrix.
# Rows represent the current state.
# Columns are the next state.
#          FOGGY  SUNNY  RAINY
FOGGY <- c( 0.5, 0.25, 0.25)
SUNNY <- c( 0.5, 0, 0.5)
RAINY <- c(0.25, 0.25, 0.5)

# Define the parameter space.
parameter <- c("Foggy", "Sunny", "Rainy")

P <- matrix(c(FOGGY, SUNNY, RAINY), nrow=3, ncol=3, byrow = TRUE,
            dimnames = list(parameter, parameter))
print(P)
```

Using what we already know, we can answer questions about weather in the future via calculation. For example, suppose today is foggy. What then is the probability that the weather will be rainy two days from now? One way is for tomorrow to be foggy, and two days from now to be rainy. Since today is foggy, the probability that tomorrow is also foggy is 0.5. And if tomorrow is foggy, then the probability that two days from now is rainy is 0.25. Thus, the probability of it being foggy then rainy is the product of these probabilities ($0.5 \times 0.25 = 0.125$). A different way is for tomorrow to be rainy (a 0.25 probability) and then rainy two days from now (a 0.5 probability), which we can determine to have a 0.125 probability of occurring

based on the product of the probabilities. The last way is for tomorrow to be sunny followed by a rainy day. Since today is foggy, the probability of tomorrow being sunny is 0.25. If tomorrow is sunny, then the probability that it will be rainy two days from now is 0.50. Hence, if today is foggy, then the probability of it being sunny then rainy is the product ($0.25 \times 0.5 = 0.125$). These three events are independent (tomorrow cannot be foggy, rainy, and sunny all at the same time), so combining these three probabilities gives a total probability of $0.125 + 0.125 + 0.125 = 0.375$ that it will be rainy two days from now. Thus, there is a 0.375 probability that the weather will be rainy two days from now given today is foggy.

The information that we already know can answer more than just the probability that the weather will be rainy two days from now given today is foggy. Similar to the example above, we could determine the probability that it will be foggy or sunny two days from now given it is foggy today. We could also determine the same weather probabilities given today was sunny or rainy instead of foggy. Additionally, this information could be used to forecast further than two days into the future. Using the steps above, try to calculate the probability that it will be foggy or sunny three days from now given that it is foggy today. You should calculate a 0.41 probability that it will be foggy and a 0.20 probability that it will be sunny three days from now.

If you continue to forecast the weather further into the future, eventually your probabilities will remain the same no matter what the weather is like today. In this example, the probabilities for the three types of weather are; $\text{Probability}(\text{rainy}) = 0.4$, $\text{Probability}(\text{sunny}) = 0.2$, and $\text{Probability}(\text{foggy}) = 0.4$. Once the probabilities remain the same no matter the initial weather (i.e., “forget the initial conditions”), the system has reached an equilibrium probability distribution of the parameters. The code block below runs this process.

In the code block below, the command `mat.or.vec()` produces a vector or matrix with a number of rows equal to the first argument and a number of columns equal to the second argument. Initially, all of the elements of this new vector or matrix have the value 0. Based on the code blocks below, what should be the dimensions of `CurrentStateProb`? Confirm your guess using the command `dim(CurrentStateProb)` after running the code block.

```
# PREDICTING THE WEATHER WITH A MARKOV CHAIN EXAMPLE
#=====
# Set the prediction length and the initial parameter value.
Prediction_Days <- 25
CurrentStateProb <- mat.or.vec(Prediction_Days, length(P[1,]))
CurrentStateProb[1,] <- c(1, 0, 0) # Today is foggy

# Run the Markov chain.
for(i in 2:Prediction_Days){
  # Current parameter value times probability
  Prob <- CurrentStateProb[i-1, ] * P
  CurrentStateProb[i, ] <- c(sum(Prob[,1]), sum(Prob[,2]), sum(Prob[,3]))
}
colnames(CurrentStateProb) <- parameter
print(CurrentStateProb)

# Print weather predictions.
print(paste("p(", parameter, ") in 1 day = ", round(CurrentStateProb[2, ], 2)))
print(paste("p(", parameter, ") in 5 days = ", round(CurrentStateProb[6, ], 2)))
print(paste("p(", parameter, ") in 24 days = ", round(CurrentStateProb[25, ], 2)))
```

Alternatively, we can achieve the same goal via MCMC. We may draw today’s weather using the probabilities for the various kinds of weather following a foggy day. The initial value is not too important at this point because we typically include a *burn-in period* to remove the effects of starting values. The burn-in period is the process of throwing away some initial portion of the Markov chain so as to remove dependence of the results on the initial conditions or in other words to forget initial conditions. This initial portion of the

Markov chain is thrown out because it is still dependent on the initial conditions. Then we iteratively draw the weather of next day based on the transition probability. After a large enough number of draws, the distribution of the samples will eventually converge to the equilibrium distribution: Probability(rainy) = 0.4, Probability(sunny) = 0.2, and Probability(foggy) = 0.4. The code block below runs this process and displays the output from both the Markov chain and MCMC. Note that the difference is subtle. The difference stems from the fact that the Markov chain calculates the exact probability while the MCMC simulates the probability.

```
# SAMPLING THE PROBABILITY DISTRIBUTION; A MARKOV CHAIN MONTE CARLO EXAMPLE
#=====
# Set the initial parameter value probability.
CurrentState <- sample(parameter, 1, prob = c(0.5, 0.25, 0.25)) # Today is foggy

# Start sampling (iteratively) from the distribution.
weather <- c()
for (i in 1:1e4){
  NextState <- sample(parameter, 1, prob = P[CurrentState,])
  weather[i] <- NextState
  CurrentState <- NextState
}

# Throw away the first 1% of the data (Burn-in)
burnin <- seq(from = 1, to = 1e4*0.1, by = 1)
weatherDraws <- weather[-burnin]

# DISPLAY
#=====
par(mfrow = c(1,2))
# Display the results from the Markov Chain
plot(1:Prediction_Days, CurrentStateProb[,1], xlab = "Days", ylab = "P(Weather)",
     ylim = c(0,1), type = "l", lwd = 2, col = "gray")
lines(1:Prediction_Days, CurrentStateProb[,2], lwd = 2, col = "gold")
lines(1:Prediction_Days, CurrentStateProb[,3], lwd = 2, col = "blue")
legend("right", c("Foggy", "Sunny", "Rainy"), lwd = 2, bty = "n", col = c("gray", "gold",
"blue"))

# Display the results from MCMC
barplot(prop.table(table(weatherDraws)), ylim = c(0,1),
        sub = "Equilibrium distribution\nof the weather", col = c("gray", "blue", "gold"))
abline(h = 0.4, lty = 2); abline(h = 0.2, lty = 2)
```

Tutorial

If you have not already done so, download the .zip file containing the scripts associated with this book from www.scrimhub.org/raes. Put the file labX_sample.R in an empty directory. Open the R script labX_sample.R and examine its contents. The tutorial example will show exactly what is coded within MCMC by having the user code MCMC rather than use a MCMC package. This code has been adapted from a class example taught by Dr. Klaus Keller (specific details are given with labX_sample.R).

In a real world application, we would have some observations of a system. In this tutorial, we will instead fit a two parameter model to some pseudo observations that we generate. We choose this approach to increase transparency of MCMC. It is also useful in understanding the assumptions made regarding the structure of the observations. In the following chapters on MCMC, we will increase the complexity, so in the end you will be able to calibrate a model to observations using MCMC.




Figure 1: The probability of the weather (rainy, sunny, or foggy) predicted over time (Markov chain process; left) and the equilibrium distribution (MCMC process; right).

The following approach and assumption is taken to approximate observations,

$$\underbrace{y_t}_{\text{observations}} = \underbrace{f(\theta, t)}_{\text{physical model}} + \underbrace{\epsilon_t}_{\text{measurement errors}},$$

$$f(\theta, t) = \alpha \times t + \beta,$$

$$\epsilon_t \sim N(0, \sigma^2),$$

$$\underbrace{\theta}_{\text{parameters}} = (\alpha, \beta),$$

where θ denotes the unknown slope (α) and intercept (β) parameters, and $f(\theta, t)$ denote the model output at parameter setting θ and time t . The observed data, y_t , is then equal to the physical model output plus measurement errors ϵ_t .

```
# Clear away any existing variables or figures and set the seed for random sampling.
rm(list = ls())
graphics.off()
set.seed(1234)


# Define the true model parameters and set up the time scale.
alpha.true <- 2 # Arbitrary choices.
beta.true <- -5
time <- 1:10

# Generate some observations with noise (measurement error) and the model.
y.true = alpha.true*time + beta.true
sigma = 1
meas_err <- rnorm(length(time), mean = 0, sd = sigma)
y = y.true + meas_err

# Plot the true model and observations.
par(mfrow = c(1,1))
plot(time, y.true, type="l", main="True observations and model",
      xlab = "Time", ylab = "Observations")
points(time, y, pch = 20)
```

Define the log-posterior

In Bayesian inference, we approximate the *posterior probability* by defining the *prior probability* and the *likelihood function*. The posterior probability is the probability that an event or observation will occur after taking into account all evidence and background information. The prior probability differs because it is our belief about the probability that an event or observation will occur, before we have taken into account new evidence. It is based on background information. In this example, you are given very little background information and thus the prior is set as an improper uniform “relatively” uninformative prior by setting the “log.prior” equal to zero. The prior is set to zero because the probability density function of a continuous uniform distribution is



where a and b are constants. We take into account new evidence using the likelihood function. The likelihood function describes the plausibility of a parameter value based on observations. Notice in the likelihood function (`log.likelihood`), we use the least squares error also known as *L2 error* rather than the least absolute deviations or *L1 error*. Using the L2 error in the likelihood function produces a “good” estimate of the mean, whereas, L1 error produces a “better” estimate of the median. Generally, we seek parameter values that maximize the likelihood function, in light of the uncertainties in both the parameters and the observations. Bayes’ theorem defines the posterior probability as proportional to the likelihood of the observations given the parameters times the prior probability of the parameters (Bayes, 1764),

$$\overbrace{p(\text{parameters} \mid \text{observations})}^{\text{Posterior}} \propto \overbrace{L(\text{observations} \mid \text{parameters})}^{\text{Likelihood}} \times \overbrace{p(\text{parameters})}^{\text{Prior}}.$$

The posterior distribution therefore summarizes information about the parameters based on the prior distribution and what the likelihood function says about more “likely” parameter values. The posterior therefore provides a probability distribution on the range of parameter values, and says which values are more probable than others.

In this analysis, you will work with the log-probability distributions for numerical stability reasons. That is, the probabilities involved may be very small, and computers may not be able to distinguish them from 0 in many cases.

```
# Define the unnormalized log posterior.
```

```
logp = function(theta){
  N = length(time)
```

```
  # Calculate model simulations.
```

```
  alpha = theta[1]; beta = theta[2]
```

```
  model = alpha*time + beta
```

```
  # Estimate the residuals (i.e., the deviation of the observations from the
  # model simulation).
```



```

resid = y - model

# Get the log of the likelihood function.
log.likelihood = -N/2*log(2*pi) - N*log(sigma) - 1/2*sum(resid^2)/sigma^2

# Use an improper uniform "relatively uninformative" prior.
log.prior = 0 # log(1)

# Bayesian updating: update the probability estimates based on the observations.
log.posterior = log.likelihood + log.prior

# Return the unnormalized log posterior value.
return(log.posterior)
}

```

Code example with Metropolis-Hastings

The algorithm used for MCMC in this tutorial is called Metropolis-Hastings where the method implemented here uses a *random walk*. A random walk is a wandering movement of an object away from where it started, following no recognizable pattern. The Metropolis-Hastings algorithm randomly samples the probability distribution and creates a Markov chain, in which the next step in the random walk only depends on the current state, and is determined by the transition probabilities and the *step size* (the proposed deviation between the current state and the next). The step size is important because if the jumps are “too small”, then the Markov chain explores the parameter space at a slower rate with a high *acceptance rate* and hence will increase the number of iterations needed for convergence. If the jumps are “too large”, then the chain could get stuck in places and possibly reject many values. The acceptance rate is the percentage of candidates that were accepted or the ratio of the number of unique values in the Markov chain to the total number of values in the Markov chain.




Figure 2: The chain of a single parameter that is poorly mixed with a small acceptance rate and a large step size (top left), poorly mixed with a large acceptance rate and small step size (top right), and well mixed and converged with a ‘good’ acceptance rate and step size (bottom left).

The acceptance rate is not to be confused with the *acceptance ratio*, which is used to decide whether to accept or reject a proposed value and is how the Markov chain is constructed. In general, the step size should be chosen so that the acceptance rate is far from 0 and far from 1 (Rosenthal, 2010). The optimal acceptance rate varies depending on the number of dimensions in the parameter space. As the number of dimensions in the parameter space increases, the optimal acceptance rate goes to about 23.4% (Roberts et al., 1997; Rosenthal, 2010). But for a single parameter, or if you implement a Metropolis-within-Gibbs (or similar) sampling approach, it is about 44% (Rosenthal, 2010). For 2 parameters, it is likely higher than 23.4%. Note that while the step size is set in this tutorial, there are adaptive algorithms for the MCMC sampler to “learn” what the step sizes ought to be, in order to achieve a desired acceptance rate. For example, there is a package called `adaptMCMC` where the MCMC function has arguments to specify a desired acceptance rate as well as the option to adapt the Metropolis sampler to achieve this desired acceptance rate (Vihola, 2011).

Constructing a Markov Chain

To construct the chains, you need to decide how long (how many iterations) to run MCMC and start with some initial parameter values, $\theta^{initial}$. In the example below, the number of iterations is set to 30,000 and the initial parameter values is randomly generated. In a “real” situation you will likely have more informed

parameter values either based on expert opinion or based on an optimization algorithm. Note that in the following chapter, you will examine how to decide how many iterations is enough.

```
# Set the number of MCMC iterations.
NI = 30000

# Start with some initial state of parameter estimates, theta^initial
alpha.init = runif(1, -50, 50) # Arbitrary choices.
beta.init = runif(1, -50, 50)
theta = c(alpha.init, beta.init)
```

Using the initial parameter values $\theta^{initial}$, you then evaluate the unnormalized posterior of that value, $p(\theta^{initial} | y)$, using the likelihood function ($L(y | \theta)$) and the prior distribution.

```
# Evaluate the unnormalized posterior of the parameter values
# P(theta^initial | y)
lp = logp(theta)
```

From here, a new parameter value, θ^{new} is proposed by being randomly drawn based on the current parameter value, transition probability ($p(\theta^{new} | \theta^{initial})$), and *step size* (the proposed deviation between the current state and the next). Unlike the weather example where we define the transition probability matrix, here, the transition probabilities are defined by a normal distribution centered at the current state with a covariance defined by the step sizes. You then evaluate the new unnormalized posterior of that value, $p(\theta^{new} | y)$, using the likelihood function. Comparing the new value to the old one, the code evaluates whether or not to accept the new value. You accept the new value with a probability based on the ratio of $p(\theta^{new} | y)/p(\theta^{initial} | y)$. In the code, this is done by first sampling the natural log of a uniformly distributed random number between one and zero. In doing so, the acceptance ratio is thus, $\min\left(\ln(1) = 1, \ln\left(p(\theta^{new} | y)/p(\theta^{initial} | y)\right) = \ln\left(p(\theta^{new} | y)\right) - \ln\left(p(\theta^{initial} | y)\right)\right)$. To put this in simpler terms, say you are at an arbitrary parameter value within the prior probability distribution. You then pick a new randomly selected parameter value. If this proposed new parameter value is “better” (higher posterior probability) than the current state, then you accept with a probability of 1; if the proposal is worse, you accept with some probability less than 1. Keeping a list of the parameter values in the model simulation will create a vector of possible values that is the Markov chain.

```
# Setup some variables and arrays to keep track of:
theta.best = theta           # the best parameter estimates
lp.max = lp                 # the maximum of the log posterior
theta.new = rep(NA,2)        # proposed new parameters (theta^new)
accepts = 0                 # how many times the proposed new parameters are accepted
mcmc.chains = array(dim=c(NI,2)) # and a chain of accepted parameters

# Set the step size for the MCMC.
step = c(0.1, 1)

# Metropolis-Hastings algorithm MCMC; the proposal distribution proposes the next
# point to which the random walk might move. For this algorithm, this proposal
# distribution is symmetric, that is P(x to x`) = P(x` to x).
for(i in 1:NI) {
  # Propose a new state (theta^new) based on the current parameter values
  # theta and the transition probability / step size
  theta.new = rnorm(2, theta, sd = step)

  # Evaluate the new unnormalized posterior of the parameter values
  # and compare the proposed value to the current state
  lp.new = logp(theta.new)
  lq = lp.new - lp
```

```

# Metropolis test; compute the acceptance ratio
# Draw some uniformly distributed random number 'lr' from [0,1];
lr = log(runif(1))

# If lr < the new proposed value, then accept the parameters setting the
# proposed new theta (theta^new) to the current state (theta).
if(lr < lq) {
  # Update the current theta and log posterior to the new state.
  theta = theta.new
  lp = lp.new

  # If this proposed new parameter value is "better" (higher posterior probability)
  # than the current state, then accept with a probability of 1. Hence, increase
  # the number of acceptations by 1.
  accepts = accepts + 1

  # Check if the current state is the best, thus far and save if so.
  if(lp > lp.max) {
    theta.best = theta
    lp.max = lp
  }
}
# Append the parameter estimate to the chain. This will generate a series of parameter
# values (theta_0, theta_1, ...).
mcmc.chains[i,] = theta
}

```

Checking the acceptance rate

After running MCMC, check the acceptance rate to determine whether the calibration appropriately explored the parameter space. That is, the chain didn't reject or accept too many values. For two parameters, the acceptance rate should be higher than 23.4%.

```

# Calculate the parameter acceptance rate; it should be higher than 23.4%.
accept.rate <- (accepts/NI) * 100
print(accept.rate)

```

Exercise

Save labX_sample.R to a new file. Now, using what you've learned throughout the previous chapters, modify the new file so that it produces a .pdf file with 6 panels that plots the results from the MCMC Metropolis-Hastings algorithm. In the following chapter, we will go in depth in analyzing the output. Specifically, you'll need to plot:

1. the true observations, true model fit, and estimated best fit over time
2. a joint scatter plot of the alpha chain `chain[,1]` versus the beta chain `chain[,2]`
3. the alpha chain and the beta chain versus the iteration number
4. histograms of the the alpha chain and the beta chain

Make sure that your plots have descriptive labels for the axes and a legend if necessary.

Questions

1. What is the purpose of defining a likelihood function?
2. In the exercise, you plotted the parameter chains, what do these chains mean/show?
3. Explain in your own words what the detailed balance equation is and why MCMC—specifically the Metropolis-Hasting algorithm—works?

References

Bayes, T. 1764. An essay toward solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society on London* 53, 370-418. Available online at <http://rstl.royalsocietypublishing.org/content/53/370.full.pdf+html>

Detailed balance. 2017. In Wikipedia, The Free Encyclopedia. Available online at https://en.wikipedia.org/w/index.php?title=Detailed_balance&oldid=810481659

Gilks, W. R. 1997. *Markov chain Monte Carlo in practice*. Chapman & Hall/CRC Press, London, UK

Grinstead, C. M. and Snell, J. L. 2006. *Introduction to Probability*. American Mathematical Society, 2006. Available online at http://www.dartmouth.edu/~chance/teaching_aids/books_articles/probability_book/amsbook.mac.pdf

Hastings, W. K. 1970. Monte Carlo sampling methods using Markov chains and their applications. *Biometrika* 57(1):97–109. doi:10.1093/biomet/57.1.97. Available online at https://www.jstor.org/stable/2334940?seq=1#page_scan_tab_contents

Howard, R. A. 2007. *Dynamic Probabilistic Systems, Volume I: Markov Models*. Dover Publications, Inc., Mineola, NY. Available online at: <https://books.google.com/books?id=DU06AwAAQBAJ>

Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., and Teller, E. Equation of state calculations by fast computing machines. *J. Chem. Phys.* 21(6), 1087-1092. Available online at <http://aip.scitation.org/doi/abs/10.1063/1.1699114>

(ML 18.6) Detailed balance (a.k.a. Reversibility). *YouTube*. Posted by mathematicalmonk, July 25, 2011. Available online at <https://www.youtube.com/watch?v=xxDkdwQdGvs>

Roberts, G. O., Gelman, A., and Gilks W. R. 1997. Weak convergence and optimal scaling of random walk Metropolis algorithms. *Ann. Appl. Prob.* 7, 110–120. Available online at http://projecteuclid.org/download/pdf_1/euclid.aop/1034625254

Rosenthal, J. S. 2010. Optimal Proposal Distributions and Adaptive MCMC. *Handbook of Markov chain Monte Carlo*. Eds., Brooks, S., Gelman, A., Jones, G. L., and Meng, X.-L. Chapman & Hall/CRC Press. Available online at <https://pdfs.semanticscholar.org/3576/ee874e983908f9214318abb8ca425316c9ed.pdf>

Ruckert, K. L., Guan, Y., Bakker, A. M. R., Forest, C. E., and Keller, K. The effects of non-stationary observation errors on semi-empirical sea-level projections. *Climatic Change* 140(3), 349-360. Available online at <http://dx.doi.org/10.1007/s10584-016-1858-z>

Schruben, L. W. 1982. Detecting initialization bias in simulation experiments. *Opns. Res.*, 30, 569-590. Available online at <http://dl.acm.org/citation.cfm?id=2754246>

Vihola, M. 2011. Robust adaptive Metropolis algorithm with coerced acceptance rate. *Statistics and Computing*. Available online at <http://www.springerlink.com/content/672270222w79h431/>

Wood, S. N. 2015. *Core Statistics, Volume 6 of Institute of Mathematical Statistics Textbooks*. Cambridge University Press. available online at <https://people.maths.bris.ac.uk/~sw15190/core-statistics.pdf>

Chapter #3: A Calibration Problem and Markov chain Monte Carlo (Kelsey L. Ruckert, Tony E. Wong, Yawen Guan, Murali Haran, and Patrick J. Applegate)

Learning objectives

After completing this exercise, you should be able to

- explain what convergence means in terms of a Markov chain
- describe several methods on how to test for convergence

Introduction

In the last chapter, we introduced Bayesian inference and Markov chain Monte Carlo (MCMC). In particular, we looked at why the Metropolis-Hastings algorithm works by satisfying detailed balance and calibrated a linear model to observations using MCMC. However, how do we know whether we can trust the results? When making inferences about probability distributions using MCMC, it is necessary to make sure that the chains (vectors of parameter values) show evidence of convergence. A converged chain is one in which the statistical properties of the chain do not change as new random values are generated. Typically, chains are not converged when they are begun and reach convergence only after many random numbers (“iterations”) have been generated.

In this exercise, you will examine multiple ways on how to test for evidence of convergence using an application of Bayesian inference with MCMC to data with non-correlated residuals. You will also see how the amount of data and how long you run the calibration impact the results and evidence of convergence. If you have not already completed the previous chapter on MCMC, go back and complete that exercise as the concepts in this chapter build off of the previous one.

Tutorial

If you have not already done so, download the .zip file containing the scripts associated with this book from www.scrimhub.org/raes. Put the file labX_sample.R in an empty directory. Open the R script labX_sample.R and examine its contents.

In this tutorial, you will estimate parameters for a linear physical model using Bayesian inference with MCMC and test for evidence of convergence. Before you proceed, first clear away any existing variables or figures and install the mcmc package and a package for testing convergence, coda.

```
# Clear away any existing variables or figures.
rm(list = ls())
graphics.off()

# Install and read in packages.
# install.packages("coda")
# install.packages("mcmc")
# install.packages("batchmeans")
library(coda)
library(mcmc)
library(batchmeans)

# Set the seed for random sampling.
# All seeds in this tutorial are arbitrary.
set.seed(1)
```




Figure 3: The original data (left) and residuals (right) as a function of time. The original data has a true value of parameter $\alpha = 0.02$, $\beta = 1$, $\sigma = 0.5$, and autocorrelation = 0.

For simplicity, the data have a linearly increasing trend over time, have independent measurement errors, and you will use the same simple model as before,

$$y_t = f(\theta, t) + \epsilon_t,$$

$$f(\theta, t) = \alpha \times t + \beta,$$

$$\epsilon_t \sim N(0, \sigma^2), \quad \theta = (\alpha, \beta),$$

where θ denotes the physical parameters, $f(\theta, t)$ is the model output, y_t is the observed data y_t , and ϵ_t are the measurement errors.

```
# Read in some observations with non-correlated measurement error
# (independent and identically distributed assumption).
data <- read.table("observations.txt", header=TRUE)
t <- data$time
observations <- data$observations

# Plot data.
par(mfrow = c(1,1))
plot(t, observations, pch = 20, xlab = "Time", ylab = "Observations")

# Set up a simple linear equation as a physical model.
model <- function(parm,t){ # Inputs are parameters and length of data
  model.p <- length(parm) # number of parameters in the physical model
  alpha <- parm[1]
  beta <- parm[2]
  y.mod <- alpha*t + beta # This linear equation represents a simple physical model
  return(list(mod.obs = y.mod, model.p = model.p))
}
```

Setting up the prior information and the likelihood function

Now that you have some data, you can obtain estimates of the parameters as well as their uncertainties. Before running the MCMC calibration, it is necessary to set up some information on what you know about

the physical model and data. For instance, you must specify initial parameter values for the Markov chains, a prior distribution for the parameters, and set up how to calculate the log posterior.

In the previous chapter, you estimated the physical model parameters (α , the slope, and β , the intercept) while setting the statistical parameter (σ , our estimate of the standard deviation) to a constant of 1. Additionally, you used initial parameter values that were randomly generated. In physical applications, it is likely better to use more informative initial parameter values that are based on, for example, a mechanistic understanding of the modeled system, previous published estimates, or by using an optimization technique. In this example, you will use the `optim` command to obtain values for α and β that minimize the root mean squared error, and use these as the Markov chain initial values, but you must, in turn, provide initial estimates for α and β to the `optim` command. These initial values can be a random guess for α and β (try $\alpha = 0.5$ and $\beta = 2$). Try using several different initial values. If the Markov chains are converged, then their distributions of parameters should converge to the same posterior distribution, regardless of the initial values specified for each chain. Note that the `optim` command will fail if there is a multimodal or complicated distribution, so try a few starting values for α and β or use the `DEoptim` command. The residuals from this optimized physical model simulation can be used to estimate a reasonable starting value for σ . Later on, you will apply the likelihood function to update these initial starting values to account for the prior parameter distributions.

```
# Sample function for calculating the root mean squared error given a set of
# parameters, a vector of time values, and a vector of observations.
```

```
fn <- function(parameters, t, obs){
  alpha <- parameters[1]
  beta <- parameters[2]
  data <- alpha*t + beta
  resid <- obs - data
  rmse <- sqrt(mean(resid^2))
  # return the root mean square error
  return(rmse)
}
```

```
# Plug in random values for the parameters.
parameter_guess <- c(0.5, 2)
```

```
# Optimize the physical model to find initial starting values for parameters.
# For optim to print more information add the arguments:
# method = "L-BFGS-B", control=list(trace=6))
result <- optim(parameter_guess, fn, gr=NULL, t, observations)
start_alpha <- result$par[1]
start_beta <- result$par[2]
parameter <- c(start_alpha, start_beta)
```

```
# Use the optimized parameters to generate a fit to the data and
# calculate the residuals.
y.obs <- model(parameter,t)
res <- observations - y.obs$mod.obs
start_sigma <- sd(res)
```

```
par(mfrow = c(1,1))
plot(res, type = "l", ylab = "Residuals", xlab = "Time")
points(res, pch = 20)
abline(h = 0, lty = 2)
```

Previously, you used priors as uniform distributions with no bounds. In physical applications, you should use a mechanistic understanding of the modeled system to set up prior parameter distributions. Assuming that there is relatively little information known about the data and physical model, how would you come

up with prior distributions? Maybe test out several values, calculate the root mean square error for several parameter sets, or look at the output from the `optim` call (try setting the method to “L-BFGS-B” and trace to 6). For this tutorial, you will set uniform prior distributions by specifying a lower and upper bound for each parameter. Note that in a physical application, using non-uniform priors may be more informative. Looking at the output from the `optim` call, it shows that α did not vary much from 0.19 and that the final root mean square error is roughly 0.46. Based on this information, you can infer that α is relatively small and therefore set the lower and upper bound to be -1 and 1. Similarly, the β parameter fluctuates between 2 and 1 in the output from the `optim` call. From this, you can set the lower and upper bound of β to be -1 and 3. Now look at the original data in Figure 2, notice how the data are closely grouped together and the residuals display deviations of less than 2 from the trend (0). Using this information, you can imply that the σ cannot be negative and that it is small, hence you can set the lower and upper bound of parameter σ to be 0 to 1. In this application, the `/sigma` is converted to variance `/sigma^2` in the likelihood function, however calibrating the variance `/sigma^2` directly is a viable option. Calibrating the variance directly should be done if using a conjugate prior such as an inverse gamma distribution.

```
# Set up priors.
bound.lower <- c(-1, -1, 0)
bound.upper <- c( 1,  3, 1)
```

In this tutorial, we provide the likelihood function, $L(y \mid \theta)$, as a separate script to be sourced. Open up `iid_obs_likelihood.R` and note the differences in how we estimate the log posterior from how it was coded in the previous chapter. Here, we split the log likelihood function, log prior, and log posterior into separate R functions. The likelihood is now the product of potentially many probabilities (between 0 and 1), so with lots of data will give underflow (Figure 1). Underflow happens when an operation performed on a value smaller than the smallest magnitude non-zero number, which is often rounded to zero. Again, we work with the log-probability distributions to solve this problem.

The log likelihood function, `log.lik`, reads in the parameters (physical and statistical) and produces model predictions. These predictions are then evaluated against the observations to estimate the residuals. The residuals along with the statistical parameter σ (the standard deviation) are used to estimate the *likelihood* of the parameter values based on the observations. Note that in the previous chapter, the log likelihood was coded as an equation,

$$L(y \mid \theta) = -\frac{N}{2} \ln(2\pi) - N \ln(\sigma) - \frac{1}{2} \frac{\sum (y - f(\theta, t))^2}{\sigma^2},$$

while here it uses the `sum` and `dnorm` commands. Both versions are the same and should produce the same value when interchanged.

In the log prior function, `log.pri`, the parameters are set as uniform distributions with an upper and lower bound. When the parameters are read in the function checks whether the parameter values are within in upper and lower bounds. If they are within the distribution, then the prior is set to 0 to represent a uniform distribution because the natural log of 1, $\ln(1)$, is 0. If the parameter values fall outside of the boundary, then the prior is set to `-Inf`. The function returns `-Inf` because there is 0 probability associated with parameters outside of the range.

Lastly, the log posterior function, `log.post`, estimates the posterior probability based on the likelihood and the prior. In this function, both the prior and likelihood functions described above are called. The parameters are first evaluated to check whether they satisfy the prior distribution. If the parameters pass with a nonzero prior probability, then the log likelihood function is called to estimate the likelihood of the parameter values based on the observations. The function then estimates the posterior probability by combining the likelihood and the prior probability. If the parameters have a prior probability of zero, then the likelihood function is not run and the posterior probability is set to `-Inf` (a probability of 0). It is important to point out that this process of not evaluating the model at parameter values outside the prior range (that will not be accepted anyway) can save valuable time. For instance, if someone was to code the log posterior function and the MCMC function by scratch, this process might be overlooked, and more expensive model (in terms of computer storage and time) can be a huge burden.

```

# Name the parameters and specify the number of physical model parameters (alpha and beta).
# sigma is a statistical parameter and will not be counted in the number.
parnames <- c("alpha", "beta", "sigma")
model.p <- 2

# Load the likelihood model for measurement errors
source("iid_obs_likelihood.R")

# Optimize the likelihood function to estimate initial starting values
p <- c(start_alpha, start_beta, start_sigma)
p0 <- c(0.3, 1, 0.6) # random guesses
p0 <- optim(p0, function(p) -log.post(p))$par
print(round(p0,4))

```

Bayesian Inference using MCMC

For the MCMC calibration, this tutorial uses the `metrop` function from the package called `mcmc`. In the console, open up the `metrop` function documentation by typing `help(metrop)`. This function runs the MCMC calibration using the Metropolis-Hasting algorithm where a random walk is implemented. Both the `metrop` function and the function we coded in the previous chapter should be the same. By using the `metrop` function, you increase the efficiency since it has already been optimized and generalized.

The `metrop` function calls for multiple input arguments including `obj`, `initial`, `nbatch`, and `scale`. The `obj` input calls for a function that estimates the unnormalized log posterior. This has been previously coded as the `log.post` function in `iid_obs_likelihood.R`, where the function estimates the posterior probability by combining the likelihood and the prior probability, as described in the previous section. The next input is `initial`, which is the initial parameter values `p0` that you calculated using the `optim` command and the likelihood function. The argument `nbatch` sets the number of iterations to run the MCMC calibration. In exercise one, you are asked to vary the number of iterations, but for the first pass set the number of iterations to 1000. Lastly, set the `scale`. The `scale` controls the step size. Remember that the step size is important because it has an impact on how well the Markov chain explores the parameter space. A step size that is “too small” will require a larger number of iterations to reach convergence, while one that is “too large” could possibly reject many acceptable values; the previous chapter displays a figure with an example depicting this issue. Here, the step sizes are set to values that are reasonable for this tutorial. How would you come up with step sizes if they were not given? Maybe test out several step size values, check out the `proposal.matrix` function in `iid_obs_likelihood.R`, or look at the MCMC function in `adaptMCMC`.

```

# Set the step size and number of iterations.
step <- c(0.001, 0.01, 0.01)
NI <- 1E3

# Run MCMC calibration.
mcmc.out <- metrop(log.post, p0, nbatch = NI, scale = step)
prechain <- mcmc.out$batch

```

Checking the acceptance rate

Check that the acceptance rate didn’t reject or accept too many values. For instance, the acceptance rate should be higher than 23.4% because it explores only three parameters.

```

# Print the acceptance rate as a percent.
acceptrate <- mcmc.out$accept * 100
cat("Accept rate =", acceptrate, "%\n")

```




Figure 4: The chain of a single parameter showing that is well mixed and converged, but the initial values differ considerably from the converged distribution. The highlighted area shows the burn-in period that is used to remove the effects of starting values.

Testing for convergence

We seek an answer to that initial question from the introduction: How do we know when the physical model is calibrated? This question is answered by diagnosing when the MCMC iterates have “converged” to sampling from the posterior distribution. There are many ways to test for convergence: visual inspection, Monte Carlo Standard Error, Gelman and Rubin diagnostic, Geweke diagnostic, Raftery and Lewis diagnostic, Heidelberg and Welch diagnostic, and many others. It is a good practice to use multiple tests to check for convergence of the Markov chains (the vector output for each parameter). The example here, will focus on four tests: (1) visual inspection, (2) Monte Carlo standard error, (3) the Heidelberg and Welch diagnostic, and (4) the Gelman and Rubin diagnostic (the potential scale reduction factor).

Before you begin testing for convergence or doing further analysis, subtract the *burn-in*. Even if the chains are converged, the first values in the sample may differ considerably from the converged distribution. Here, we throw away the first 1% of the sample for the *burn-in period*. Even though a 1% burn-in is used here, other applications may require different lengths or may not need a burn-in period.

Identify the burn-in period and subtract it from the chains.

```
burnin <- seq(1, 0.01*NI, 1)
mcmc.chains <- prechain[-burnin, ]
```

Visual inspection

One way to check for convergence is by looking at how the chain jumps around in the parameter space or how well it mixes. You can visualize this using a plot of the iteration number versus the parameter value (a trace plot, also known as a history plot). If the Markov chain becomes stuck in certain areas of the parameter space or if the chain is taking a long time to explore the parameter space, then the chain has not converged and indicates that either the chain needs to be run longer or there is some other kind of issue.

Check #1: Trace Plots:

```
par(mfrow = c(2,2))
for(i in 1:3){
  plot(mcmc.chains[,i], type="l", main = "",
       ylab = paste('Parameter = ', parnames[i], sep = ' '), xlab = "Number of Runs")
}
```

```
}
```

Monte Carlo standard error

The point of using MCMC is to approximate expectations. These approximations are not exact, but are off by some amount known as the *Monte Carlo standard error* (MCSE). A way to calculate the MCSE is through the *batch means method*. The batch means method estimates uncertainty by calculating means and standard errors (as square root of sample variances) of *batch means* of the markov chain (Jones et al. 2006; Flegal et al., 2008). To perform this method, suppose you have a markov chain $X = \{X_1, X_2, X_3, \dots\}$. The Monte carlo estimate or the sample mean is

$$\bar{g}_n := \frac{1}{n} \sum_{i=1}^n \bar{X}_i \quad \text{as } n \rightarrow \infty,$$

where n is the number of iterations in the markov chain and \bar{X}_i is a sample mean for a batch. In batch means, the output is broken into a number of blocks or batches of b size. By dividing the simulation into batches, the batch means estimate of σ_g^2 becomes

$$\hat{\sigma}_g^2 = \frac{b}{a-1} \sum_{j=1}^a (\bar{Y}_j - \bar{g}_n)^2$$

$$\bar{Y}_j := \frac{1}{b} \sum_{i=(j-1)b+1}^a \bar{X}_i \quad \text{for } j = 1, \dots, a.$$

The σ is denoted as $\hat{\sigma}$ because it is unknown and is estimated from the data. Using consistent batch means, you can estimate the MCSE of \bar{g}_n as

$$\text{MCSE} = \frac{\hat{\sigma}_g}{\sqrt{n}}$$

This method is performed in R via the `batchmeans` package. The command `bm()` performs consistent batch means estimation on a Markov chain returning the mean and the MCSE. Below, `bmmat()` is used because it performs consistent batch means estimation on a matrix of Markov chains and in this case you have three chains.

```
## Check #2: Monte Carlo Standard Error:
# est: approximation of the mean
# se: estimates the MCMC standard error
bm_est <- bmmat(mcmc.chains)
print(bm_est)
```

In theory, if you run the chain infinitely, then the standard errors of the approximations will be virtually 0. In practice, since the chain is not run infinitely, you can calculate the standard error of the approximations and decide if it is “small enough” to stop the chain. You can calculate if the estimate is small enough by assessing whether we trust the amount of significant figures in the standard error. The first step after calculating the MCSE and the Monte Carlo estimate is to estimate the z-statistic,

$$z = \frac{(\mu_X - \bar{g}_n)}{\frac{\hat{\sigma}_g}{\sqrt{n}}},$$

where μ_X is the mean of the chain, \bar{g}_n is the Monte Carlo estimate, and $\frac{\hat{\sigma}_g}{\sqrt{n}}$ is the MCSE. Here, we use the z-distribution instead of the t-distribution because a t-distribution with a degrees of freedom (\sqrt{n}) equal to or greater than 30 is very close to the z-distribution. Typically, MCMC is run with a large number of iterations (tens of thousands), which exceeds a degrees of freedom of 30. Unless the sample is highly correlated or has a

low degrees of freedom, the z-distribution can be used. Next, is to calculate the half-width h_α of the interval by multiplying the z-statistic and the MCSE,

$$h_\alpha = z \left(\frac{\hat{\sigma}_g}{\sqrt{n}} \right).$$

If the lower and upper bound of the confidence intervals ($\bar{g}_n \pm h_\alpha = \bar{g}_n \pm z \left(\frac{\hat{\sigma}_g}{\sqrt{n}} \right)$ for $n \geq 30$) can be rounded to equal the Monte Carlo estimate, then the significant figure in the MCSE can be trusted. If the error is small enough, then no more samples are needed and the chain is considered converged. Additionally, reporting the MCSE provides a measure of the accuracy and quality of the estimates (Flegal et al., 2008).

```
# Evaluate the number of significant figures
z <-
half_width <- rep(NA, length(parnames))
interval <- matrix(data = NA, nrow = 3, ncol = 2,
                    dimnames = list(c(1:3), c("lower_bound", "upper_bound")))
for(i in 1:length(parnames)){
  z[i] <- (mean(mcmc.chains[,i]) - bm_est[i, "est"])/bm_est[i, "se"]
  half_width[i] <- z[i] * bm_est[i, "se"]
  interval[i,1] <- bm_est[i, "est"] - half_width[i]
  interval[i,2] <- bm_est[i, "est"] + half_width[i]
}
print(interval)
```

Heidelberger and Welch diagnostic

The Heidelberger and Welch diagnostic tests whether each parameter chain has stabilized. This test is split into two parts. The first part of the test examines whether the chain comes from a *stationary distribution*. A stationary distribution means that no matter the starting state, the distribution is unchanged and stable over time. If the chain does not come from a stationary distribution, then the first 10% of the chain will be thrown away and the test runs again. The test will continue to repeat this process until either the chain passes the test or half of the chain has been thrown away and still failed. If the first part is passed, then the test will move on to the second part using the portion of the chain that passed the stationarity test. In part two, a 95% confidence interval is calculated for the mean of the remaining chain(s). *Half of the width* of the 95% confidence interval is compared to the mean. If the ratio between the half-width number and the mean is less than the user-defined tolerance value, ϵ , then the chain passes the test. If the test fails in either part, this suggests that the chains have not converged and need to be run longer. In R, the Heidelberger and Welch diagnostic is carried out in the `heidel.diag` function. Type `help(heidel.diag)` and check out the function arguments.

```
## Check #3: Heidelberger and Welch's convergence diagnostic:
heidel.diag(mcmc.chains, eps = 0.1, pvalue = 0.05)
```

Gelman and Rubin diagnostic

The Gelman and Rubin diagnostic evaluates the *potential scale reduction factor*. The potential scale reduction factor monitors the variance within the chain and compares it against the variance between the chains and if they look similar, then this is evidence for convergence. It is necessary to estimate the “between-chain variance”, so you need to run at least one additional MCMC calibration to estimate the potential scale reduction factor. Each additional MCMC calibration needs to be run with a different random seed and a different set of starting values. All seeds in this tutorial are arbitrary and the different initial values are to ensure that posteriors are independent from the initial values.

In this case, run the MCMC calibration three more times. Each set of chains must be converted into MCMC objects and combined into a list in order to use the R routine `gelman.diag` to calculate these diagnostics.




Figure 5: The potential scale reduction factors as a function of the last iteration in the chain. This parameter chain (left) flattens out to one and is hence converged. On the right, the parameter chain has not converged.

The Gelman and Rubin diagnostic then estimates the median and 97.5% quantile of the potential scale reduction factor. We adopt the standard that the chains are considered to be converged if the point estimate and the upper limit potential scale reduction factor is no larger than 1.1. The results can also be visualized using the `gelman.plot()` function. Run the MCMC calibration longer, if the test fails.

```
## Check #4: Gelman and Rubin's convergence diagnostic:
set.seed(111)
p0 <- c(0.05, 1.5, 0.6) # Arbitrary choice.
mcmc.out2 <- metrop(log.post, p0, nbatch=NI, scale=step)
prechain2 <- mcmc.out2$batch

set.seed(1708)
p0 <- c(0.1, 0.9, 0.3) # Arbitrary choice.
mcmc.out3 <- metrop(log.post, p0, nbatch=NI, scale=step)
prechain3 <- mcmc.out3$batch

set.seed(1234)
p0 <- c(0.3, 1.1, 0.5) # Arbitrary choice.
mcmc.out4 <- metrop(log.post, p0, nbatch=NI, scale=step)
prechain4 <- mcmc.out4$batch

# The burn-in has already been subtracted from the first chain.
# Thus, the burn-in only needs to be subtracted from the three other
# chains at this point.
mcmc1 <- as.mcmc(mcmc.chains)
mcmc2 <- as.mcmc(prechain2[-burnin, ])
mcmc3 <- as.mcmc(prechain3[-burnin, ])
mcmc4 <- as.mcmc(prechain4[-burnin, ])

set.seed(1) # revert back to original seed
mcmc_chain_list <- mcmc.list(list(mcmc1, mcmc2, mcmc3, mcmc4))
gelman.diag(mcmc_chain_list)
gelman.plot(mcmc_chain_list)
```

Analyzing MCMC output

Once there is evidence that the MCMC chains are converged, the output can be analyzed to obtain information about the physical and statistical model parameters such as how confident we are about the estimates of uncertain values. For instance, you can look at the probability density of each parameter using the `density()` function. Estimates of the mean, mode, median, and credible intervals can be determined as well. The *equal-tail* credible interval can be computed using the `quantile()` function or the *highest posterior density* credible interval can be estimated using the `HPDinterval()` function. The equal-tail credible interval excludes the same percentage from each tail of the distribution. By definition, a 90% equal-tail credible interval would exclude 5% from the left and right tail so it would be calculated with the 5% estimate and the 95% estimate. The highest posterior density credible interval differs by estimating the shortest interval in the parameter space containing the specified percentage (e.g. 90%) of the posterior probability.

```
# Calculate the 90% highest posterior density CI.
# HPDinterval() requires an mcmc object; this was done in the code block above.
hpd = HPDinterval(mcmc1, prob = 0.90)

# Create density plot of each parameter.
par(mfrow = c(2,2))
for(i in 1:3){
  # Create density plot.
  p.dens = density(mcmc.chains[,i])
  plot(p.dens, xlab = paste('Parameter =', ' ', parnames[i], sep = ' '), main="")

  # Add mean estimate.
  abline(v = bm(mcmc.chains[,i])$est, lwd = 2)

  # Add 90% equal-tail CI.
  CI = quantile(mcmc.chains[,i], prob = c(0.05, 0.95))
  lines(x = CI, y = rep(0, 2), lwd = 2)
  points(x = CI, y = rep(0, 2), pch = 16)

  # Add 90% highest posterior density CI.
  lines(x = hpd[i, ], y = rep(mean(p.dens$y), 2), lwd = 2, col = "red")
  points(x = hpd[i, ], y = rep(mean(p.dens$y), 2), pch = 16, col = "red")
}
```

Exercise

Before you proceed, **make sure that you have a working version of the code blocks above.** Open your script in RStudio and `source()` it. Examine the convergence tests. You should see that the chains are visually poorly mixed and the Gelman and Rubin diagnostic failed, but the Heidelberger and Welch diagnostic passed. If not, check your script against the instructions above. Once you are satisfied that your script is working properly, save the file.

Part 1: Converging the MCMC results. Modify your new script by changing the number of iterations *NI* in block 5. The goal is to get the chains to converge. Each time you modify the number of iterations save each chain. Remember to store the parameter results (the chains) under different names each time you change the iteration number so the values are not rewritten, perhaps *chains.1thous*, *chains.10thous*, *chains.100thous*, and so on. Repeat this process until you have evidence that the MCMC chains are converged.

Your script should generate the following plots:

1. a three-panel figure with a trace plot of each parameter showing that the chains are well-mixed and converged.




Figure 6: Probability density plot of each parameter along with the mean as a vertical line, the 90% equal-tail CI as a black horizontal line, and the 90% highest posterior density CI as a red horizontal line.

2. a three-panel figure of the potential scale reduction factor for each parameter showing that the chains have converged.
3. plots of the posterior probability densities: density plots of the α , β , and σ values. Each plot should include the true parameter value from the original data as a vertical line and the densities produced from each time you changed the iteration number. Make sure to label the axes of your plots in a sensible way and include a legend.

Part 2: Varying length of data. Suppose you did not have 200 years worth of data and instead only had 20 years of data. Modify your new script by changing *num.of.obs* to 20, but keep the iteration number at the number when the results converged with 200 data points. If the results are not converged, then increase the number of iterations used for the calibration. Once satisfied, store the parameter results. Repeat this process with 50 and 100 data points and store those results.

Your script should generate the following plots:

1. plots of the posterior probability density functions of the parameter values. Each plot should include the true parameter value from the original data as a vertical line and the densities produced from each time you changed the number of data points. Make sure to label the axes of your plots and include a legend.

Questions

1. What is the minimum chain length needed in order for the MCMC calibration in this problem to converge (to the nearest 100 thousand)?
2. Compare the mean and median from each converged chain. Are the estimates close to the true parameters ($\alpha = 0.02$, $\beta = 1$, $\sigma = 0.5$)?
3. Compare the resulting parameter densities produced from different numbers of iterations. Does it matter if the calibration has converged and why?
4. How do the results change by having less data? Does the number of iterations to run until convergence change? Does the uncertainty surrounding each parameter change?
5. When you first source the code above both the trace plots and Gelman and Rubin diagnostic fail, yet the Heidelberger and Welch diagnostic suggest the chains have converged. Why does this happen and how would you rectify this issue?

Appendix

These questions are intended for students with advanced backgrounds in Statistics or the Earth Sciences and R programming.

In the tutorial, you used uniform priors. In other applications non-uniform priors maybe more appropriate. How would one go about setting up a non-uniform prior?

When you incorporate uncertainty, the range is called a confidence interval in frequentist statistics and is a credible interval in Bayesian statistics. The difference between a frequentist 90% (5-95%) confidence interval and a Bayesian 90% (5-95%) credible interval is that the confidence interval states if you do this many times, the estimate will be in there 90% of the time, whereas the Bayesian credible interval states that the physical model, data, and the statistical approach indicate that there is a 90% chance of capturing the true value. Does the credible interval capture the true values?

In this chapter, you fit a linear model to data that have non-correlated (independent and identically distributed) residuals. In the Earth sciences, the data being evaluated may have correlated residuals. How could one test whether the residuals are correlated? If the data have correlated residuals, how might one modify the MCMC calibration to account for this correlation?

References

- Bayes, T. 1764. An essay toward solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society on London* 53, 370-418. Available online at <http://rstl.royalsocietypublishing.org/content/53/370.full.pdf+html>
- Brooks, S. P. and Gelman, A. 1998. General methods for monitoring convergence of iterative simulations. *Journal of Computational and Graphical Statistics*, 7, 434-455. Available online at <http://www.tandfonline.com/doi/abs/10.1080/10618600.1998.10474787>
- Flegal, J. M., Haran, M., and Jones, G. L. 2008. Markov chain Monte Carlo: Can we trust the third significant figure? *Statistical Science*, 23, 250-260. Available online at http://sites.stat.psu.edu/~mharan/mcse_rev2.pdf
- Gelman, A. and Rubin, D. B. 1992. Inference from iterative simulation using multiple sequences. *Statistical Science*, 7, 457-511. Available online at <https://projecteuclid.org/euclid.ss/1177011136>
- Gilks, W. R. 1997. *Markov chain Monte Carlo in practice*. Chapman & Hall/CRC Press, London, UK
- Grinstead, C. M. and Snell, J. L. 2006. *Introduction to Probability*. American Mathematical Society, 2006. Available online at http://www.dartmouth.edu/~chance/teaching_aids/books_articles/probability_book/amsbook.mac.pdf
- Hastings, W. K. 1970. Monte Carlo sampling methods using Markov chains and their applications. *Biometrika* 57(1):97-109. doi:10.1093/biomet/57.1.97. Available online at https://www.jstor.org/stable/2334940?seq=1#page_scan_tab_contents
- Heidelberger, P. and Welch, P. D. 1981. A spectral method for confidence interval generation and run length control in simulations. *Comm. ACM*. 24, 233-245. Available online at <http://dl.acm.org/citation.cfm?id=358630>
- Heidelberger, P. and Welch, P. D. 1983. Simulation run length control in the presence of an initial transient. *Opns Res.* 31, 1109-1144. Available online at <http://dl.acm.org/citation.cfm?id=2771121>
- Howard, R. A. 2007. *Dynamic Probabilistic Systems, Volume I: Markov Models*. Dover Publications, Inc., Mineola, NY. Available online at <https://books.google.com/books?id=DU06AwAAQBAJ>
- Jones, G. L., Haran, M., Caffo, B. S. and Neath, R. 2006. Fixed-width output analysis for Markov chain Monte Carlo. *Journal of the American Statistical Association*, 101, 1537-1547. Available online at https://www.jstor.org/stable/27639771?seq=1#page_scan_tab_contents
- Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., and Teller, E. Equation of state calculations by fast computing machines. *J. Chem. Phys.* 21(6), 1087-1092. Available online at <http://aip.scitation.org/doi/abs/10.1063/1.1699114>
- Roberts, G. O., Gelman, A., and Gilks W. R. 1997. Weak convergence and optimal scaling of random walk Metropolis algorithms. *Ann. Appl. Prob.* 7, 110-120. Available online at http://projecteuclid.org/download/pdf_1/euclid.aoap/1034625254
- Rosenthal, J. S. 2010. "Optimal Proposal Distributions and Adaptive MCMC." *Handbook of Markov chain Monte Carlo*. Eds., Brooks, S., Gelman, A., Jones, G. L., and Meng, X.-L. Chapman & Hall/CRC Press. Available online at <https://pdfs.semanticscholar.org/3576/ee874e983908f9214318abb8ca425316c9ed.pdf>
- Ruckert, K. L., Guan, Y., Bakker, A. M. R., Forest, C. E., and Keller, K. The effects of non-stationary observation errors on semi-empirical sea-level projections. *Climatic Change* 140(3), 349-360. Available online at <http://dx.doi.org/10.1007/s10584-016-1858-z>
- Schruben, L. W. 1982. Detecting initialization bias in simulation experiments. *Opns. Res.*, 30, 569-590. Available online at <http://dl.acm.org/citation.cfm?id=2754246>

Chapter #4: Applying Markov chain Monte Carlo to sea-level data (Kelsey L. Ruckert, Tony E. Wong, Yawen Guan, and Murali Haran)

Learning objectives

After completing this exercise, you should be able to:

- perform a Markov chain Monte Carlo analysis that accounts for autocorrelation and heteroskedasticity
- apply this approach to “real” data (sea level) and use a “real” physical model (Rahmstorf 2007)
- explain why accounting for the error structure is important

Introduction

Suppose you have a physical model that characterizes the relationships within a system. However, this model is not perfect. There is an additional element of uncertainty. Adding a residual term is intended to account for this inaccuracy or uncertainty by accounting for model structural uncertainty, parametric uncertainty, physical process variability, and observational uncertainty not represented in the physical model. For example, sea-level data may have years when measurements are less or more certain simply because the instruments making these measurements may periodically break down, technological advancements increase the accuracy of these measurements over time, and more measurements become available. Further, deviations from the main trend of the data (residuals) could be dependent on previous observations. Testing the data for interdependent (autocorrelated) residuals and time-varying observation error (heteroskedasticity) is important because these can impact what assumptions you make about the data and the methods you choose to fit a physical model to the data.

In the previous chapters, you performed a Markov chain Monte Carlo (MCMC) calibration for a linear physical model, using data generated with non-correlated residuals. However, in many situations, observations in the Earth sciences (e.g., sea-level data) have a different and often more complicated error structure. Hence, a different approach for performing a MCMC analysis must be applied if you want to analyze data with a more complicated error structure. For instance, sea-level data is heteroskedastic, meaning that the error of the measurements varies over time. Additionally, sea-level residuals are autocorrelated, meaning that the current value is dependent upon previous values. In order to perform a MCMC analysis using sea-level data, it is necessary to check for evidence of these properties in the data. If evidence exists for this error structure, then it is important to account for autocorrelation and heteroskedasticity in the MCMC analysis that you perform.

How do we account for the error structure?

The error structure is defined within the residual term. It is important to distinguish the difference between residuals and errors. An error is the difference of the observed value from the “unobservable” true value. On the other hand, a residual is the deviation of an observation from its estimated value or modeled value,

$$\overbrace{R_t}^{\text{residual}} = \overbrace{y_t}^{\text{observation}} - \overbrace{f(\theta, t)}^{\text{physical model}},$$

where t is time and θ is the physical model parameters. The point of the residuals is to describe recognized amounts of unexplained variation in observations that are not always represented in a physical model. This variation can include a variety of uncertainties including model structure uncertainty, parametric uncertainty, physical process variability, and observational uncertainty. In this case, we approximate the residuals as the sum of the model error ω_t and observational/ measurement errors ϵ_t ,

$$R_t = \omega_t + \epsilon_t.$$

The model error is a time series from a multivariate normal distribution. This multivariate normal distribution is described with the variance σ_{AR1}^2 and the correlation structure given by the first-order autoregressive parameter ρ ,

$$\omega_t = \rho \times \omega_{t-1} + \underbrace{\delta_t}_{\text{white noise}},$$

$$\omega_0 \sim N(0, \frac{\sigma_{AR1}^2}{1 - \rho^2}), \quad \delta_t \sim N(0, \sigma_{AR1}^2), \quad \epsilon_t \sim N(0, \sigma_{\epsilon,t}^2).$$

This model is also recognized as a lag-1 autoregressive process or *AR1*. An autoregressive model is a representation of random time-varying processes in nature. In an autoregressive model, the output depends on its previous values and white noise. White noise is a sequence of continuously uncorrelated random variables with a zero mean and a constant variance σ_{AR1}^2 . For AR1, only the previous term in the process, the autocorrelation coefficient, and the white noise are used to create the output. If the autocorrelation runs deeper than just an AR1 process, then it may be necessary to fit more sophisticated models (e.g., lag-2 autoregressive process or other types of autoregressive integrated moving average models).

When you account for an AR1 process, you not only modify the equation of how to approximate the residuals, but also modify the likelihood function (check out `Rahm_obs_likelihoood_AR.R` and compare it to `iid_obs_likelihoood_AR.R`). For instance, the modified likelihood function accounts for the observational errors ϵ_t and the autocorrelation coefficient ρ so that the likelihood function is then,

$$L(y | \theta) = \left(\frac{1}{\sqrt{2\pi}} \right)^N \times |\Sigma|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} [y - f(\theta, t)]^T \Sigma^{-1} [y - f(\theta, t)] \right\},$$

where Σ is the sum of the variance of the AR1 process and the observation errors. The variance of the AR1 process and the variance of the observation errors are both diagonal matrices of size $N \times N$ where the diagonal has the variance/ error element (see the supplementary material of Ruckert et al. 2017 for details). Additionally, ρ is added as an uncertain parameter in the MCMC analysis.

In the previous chapter, we make a key assumption in the MCMC framework that the residuals are normally distributed with a mean of 0. However, in the example of calibrating changes in sea level, we fail this key assumption. Instead, the AR1 process is a Gaussian process, which is *white* (statistically uncorrelated), thereby the residuals satisfy the MCMC assumptions. Hence, in this example, we fit the AR1 error model and use the AR1 residuals. This effect of going from non-normally distributed and potentially non-mean 0 residuals to those that are normally distributed with mean 0 is called “whitening” the residuals, referring to the fact that often the signal strength of *white noise* is assumed to be normally distributed. There are actually many situations where we do not know the most appropriate error model to use, so there are many different ways we can “whiten” the residuals ($R_{t+1} - \rho R_t$) in order to satisfy the assumptions underlying the MCMC framework. Once we whiten the residuals we can then estimate the log likelihood as the sum of the normal density of the non-correlated (whitened) residuals with a mean of 0 and a standard deviation of the square root of the variance σ^2 plus the observation errors squared.

```
# Estimate the log likelihood of the AR1 process
logl.ar1 <- function(r, sigma1, rho1, eps1 = y.meas.err){
  n <- length(r) # r is the residuals

  logl <- 0
  if(n>1) {
    # Whiten the residuals (w)
    w <- r[2:n] - rho1*r[1:(n-1)]
    logl <- logl + sum(dnorm(w, sd=sqrt((sigma1)^2 + (eps1[c(-1)])^2), log=TRUE))
  }
  return(logl)
}
```

Here, we treat the estimation of the log prior distribution function and the log posterior distribution function the same as in the previous chapter. In this tutorial, you will use uniform distributions as priors for the

physical and statistical parameters. Hence, the log prior function evaluates whether the parameter values are within the boundary and if so returns zero to represent a uniform distribution. As before, the log posterior function estimates the posterior distribution based on the log likelihood and log prior (posterior \propto likelihood \times prior).

Tutorial

If you have not already done so, download the .zip file containing the scripts associated with this book from www.scrimhub.org/raes. Put the file `labX_sample.R` in an empty directory. Open the R script `labX_sample.R` and examine its contents. This chapter builds off of the previous chapters about MCMC, so it is highly recommended that you complete those chapters before you proceed. In this exercise, you will calibrate the Rahmstorf (2007) model to sea-level data while accounting for the autocorrelated residuals and the heteroskedastic observation errors.

Before you proceed, first clear away any existing variables or figures and install packages. Install the `adaptMCMC` package. The `compiler` package should already exist in the system library and therefore does not need to be installed.

In the code block below, the command `enableJIT()` is used to increase the efficiency of R by enabling or disabling “just-in-time” compilation. Enabling “just-in-time” compilation causes the closures to be compiled before they are first used (`enableJIT(1)`), closures to be compiled before they are duplicated (`enableJIT(2)`), and all `for()`, `while()`, and `repeat()` loops to be compiled before they are executed (`enableJIT(3)`). “Just-in-time” compilation is turned off if the argument is set to 0. Test out the speed gain by tracking the time it takes to run a set of code with `proc.time()`.

```
# Clear away any existing variables or figures.
rm(list = ls())
graphics.off()

# Install and read in packages.
# install.packages("adaptMCMC")
library(adaptMCMC)
library(compiler)
enableJIT(3)

# Set the seed.
set.seed(111)

# Make directory for this lab
dir.create("lab_X")
dir.path <- paste(getwd(), "/lab_X/", sep="")
scrim.git.path <- "https://raw.githubusercontent.com/scrin-network/"
```

Next, download and read in the sea-level observations, temperature observations, and the Rahmstorf (2007) global sea-level model. Specific details on these data sources can be found within the code block below and in previous chapters.

```
# Global mean sea level from Church and White (2011)
# year, time in years from 1880 to 2013 at the half year mark
# slr, global mean sea level in mm
# err.obs, global mean sea level measurement error in mm
url <- paste(scrim.git.path, "BRICK/master/data/GMSL_ChurchWhite2011_yr_2015.txt", sep="")
download.file(url, paste(dir.path, "GMSL_ChurchWhite2011_yr_2015.txt", sep=""))
church.data <- read.table("lab_X/GMSL_ChurchWhite2011_yr_2015.txt")
year <- church.data[,1] - 0.5
```

```

slr <- church.data[,2]
err.obs <- church.data[,3]

# Calculate sea level +/- observation errors.
err_pos <- slr + err.obs
err_neg <- slr - err.obs

# Read in the information from the Smith (2008) temperature file.
# project.years, time in years from 1880 to 2100.
# hist.temp, historical global mean temperatures from 1880 to 2013 in C
# rcp85, merged historical + rcp 8.5 temperatures from 1880 to 2100 in C
url <- paste(scrim.git.path,
             "Ruckertetal_SLR2016/master/RFILES/Data/NOAA_IPCC_RCPtempsscenarios.csv",
             sep="")
download.file(url, paste(dir.path, "NOAA_IPCC_RCPtempsscenarios.csv", sep=""))
temp.data <- read.csv("lab_X/NOAA_IPCC_RCPtempsscenarios.csv")
project.years <- temp.data[1:match(2100, temp.data[,1]), 1]
hist.temp <- temp.data[1:match(2013, temp.data[,1]), 2]
rcp85 <- temp.data[1:match(2100, temp.data[,1]), 4]

# The sea level and temperature data have a yearly timestep.
tstep = 1

# Load the sea-level model (Rahmstorf, 2007)
url <- paste(scrim.git.path, "BRICK/master/R/gmsl_r07.R", sep="")
download.file(url, paste(dir.path, "gmsl_r07.R", sep=""))
source("lab_X/gmsl_r07.R")

```

Characterizing the error structure

In the introduction, we discuss how scientists often must make assumptions about the data and the error structure of the data. In practice, how do we know what assumptions to make? To answer this, you can examine the observations and inform assumptions based on what information and details are known about the observations. For example, the sea-level observations `slr` start from the year 1880 and you can assume that since 1880 the measurements of sea level have become more accurate and more available. Since the data from Church and White (2011) includes error estimates, you can test that assumption of whether the data has time-varying observation error (heteroskedasticity) by plotting the errors over time. The resulting plot shows that the errors indeed vary over time.

```

par(mfrow = c(1,1))
plot(year, err.obs, type = "l", ylab = "Observation errors [mm]", xlab = "Time")
points(year, err.obs, pch = 20)

```

Additionally, you can assume that changes in sea level this year will impact what sea level is like next year. Hence, this implies that the residuals are autocorrelated. Before testing for autocorrelation, you must calculate the residuals. Using an optimization procedure (`optim` or `DEoptim`), you can estimate the parameter values. The model parameters include the sensitivity of sea level to changes in temperature, α ; the temperature when then sea-level anomaly is zero, T_{eq} ; and the initial sea-level anomaly value (the value in 1880), SL_0 . These values can be plugged into the model to generate a “best fit” simulation and be later used as initial parameter estimates. Using the resulting “best-fit” simulation, you estimate the residuals by subtracting the model simulation from the observations. Taking the standard deviation of the residuals can serve as the initial estimate of σ later on.

```

# Generate fit to sea-level observations using General-purpose Optimization.
# by calculating the root mean squared error.
fn <- function(pars, tstep, Temp, obs){
  a <- pars[1]
  Teq <- pars[2]
  SL0 <- pars[3]

  np <- length(Temp)
  gmsl <- rep(NA, np)
  gmsl[1] <- SL0

  # Use Forward Euler to estimate sea level over time.
  for (i in 2:np) {
    gmsl[i] <- gmsl[i-1] + tstep * a * ( Temp[i-1]-Teq )
  }
  resid <- obs - gmsl
  # Return the root mean square error
  rmse <- sqrt(mean(resid^2))
  return(rmse)
}

# Plug in random values for the parameters.
parameter_guess <- c(3.4, -0.5, slr[1])

# Optimize the model to find initial starting values for parameters.
result <- optim(parameter_guess, fn, gr=NULL, tstep, Temp = hist.temp, obs = slr)
start_alpha <- result$par[1]
start_Teq <- result$par[2]
start_SL0 <- result$par[3]

# Make a plot of observations and fit.
slr.est <- gmsl_r07(a = start_alpha, Teq = start_Teq, SL0 = start_SL0, tstep,
                  Tg = hist.temp)
plot(year, slr, pch = 20, ylab = "Sea-level Anomaly [mm]", xlab = "Year")
lines(year, slr.est, col = "blue", lwd = 2)

# Calculate residuals and initial sigma value.
res <- slr - slr.est
start_sigma <- sd(res)

```

You can use the resulting residuals to test for autocorrelation with the `acf()` function. “Autocorrelation” refers to the correlation between two points within a time series, where the time difference between them is called the “lag”. In the autocorrelation plot, the correlation is shown at each time lag. If the correlation exceeds the dashed blue lines (a 95% confidence interval), then the correlation of the data at that lag is considered to be “statistically significant”. Values below the dashed lines are considered insignificant correlations. The correlation at lag 0 should always equal 1 because each observation should be 100% correlated to itself. The resulting plot shows that the residuals are autocorrelated with the highest correlation at the lag-1 autoregressive process or *AR1*. Accounting for the *AR1* process requires adding an uncertain statistical parameter, ρ . ρ is the first-order or lag-1 autocorrelation coefficient. For later use, you can save `rho[2]` as the initial ρ estimate to set up the MCMC calibration.

```

# Apply the auto-correlation function to determine correlation coefficients.
rho <- rep(NA,3)
ac <- acf(res, lag.max = 5, plot = TRUE, main = "", xlab = "Time lag",

```

```

      ylab = "Autocorrelation function")
rho[1] <- ac$acf[1]
rho[2] <- ac$acf[2]
rho[3] <- ac$acf[3]
rho[4] <- ac$acf[4]
rho[5] <- ac$acf[5]
start_rho <- rho[2]

```

Setting up the prior information

In the previous blocks, you read in the data, learned about what to account for in the error structure, and estimated some good initial parameter estimates. The next goal is to learn about the distribution of the parameters (α , T_{eq} , SL_0 , σ , and ρ) given the observed data. First, you will specify a prior distribution for each model parameter and statistical parameter. Like the previous MCMC chapter, you will use uniform prior distributions. For the model parameters, you will use the prior distributions based on Ruckert et al. (2017; see supplementary material). Hence, the prior for α will be set as 0 to 20 [$mm \cdot yr^{-1} \cdot ^\circ C^{-1}$], -3 to 2 [$^\circ C$] for T_{eq} , and as the sea-level observation [mm] in 1880 plus and minus the observation error in 1880 (1880 is the year of the first sea-level value) for SL_0 . Also, set σ as 0 to 10 [mm] because in this example the σ cannot be negative and that it is small based on observing the residuals. ρ is set to have a prior range of -0.99 to +0.99. This is because ρ is a percent represented as a decimal and therefore cannot have a positive or negative correlation greater than 100%. As an added challenge (recommended in the Appendix), test out different nonuniform prior distributions for the model parameters.

The difference with the approach shown in the previous MCMC chapter is that the data have residuals that are correlated and show heteroskedasticity. To account for the differences in the error structure, we update the likelihood function to add ρ as an uncertain parameter and to set the measurement errors equal to the known observational errors, `err.obs` (as stated in the introduction).

```

# Set up MCMC calibration.
# Set up priors.
bound.lower <- c( 0, -3, err_neg[1], 0, -0.99)
bound.upper <- c(20,  2, err_pos[1], 10,  0.99)

# Name the parameters and specify the number of physical model parameters.
# Sigma and rho are statistical parameters and will not be counted in the number.
parnames <- c("alpha", "Teq", "SL0", "sigma", "rho")
model.p <- 3

# Set the measurement errors.
y.meas.err <- err.obs

# Load the likelihood model accounting for the error structure.
source("Rahm_obs_likelihood_AR.R")

# Setup initial parameters.
p <- c(start_alpha, start_Teq, start_SL0, start_sigma, start_rho)

# Set p0 to a vector of values to be used in optimizing the log-posterior function.
# We optimize the negative log-posterior function because "optim" only knows how to
# minimize a function, while we would like the log-posterior to be maximized.
p0 <- c(3.4, -0.5, slr[1], 6, 0.5)
p0 <- optim(p0, function(p) -log.post(p))$par
print(round(p0, 4))

```


Running MCMC

For this chapter, you will use a more sophisticated MCMC algorithm with an adaptive algorithm for the MCMC sampler (Vihola, 2011). This adaptive MCMC uses Metropolis-Hastings updates with joint multivariate normal proposal, which tunes the covariance matrix for these proposals to “learn” what the step sizes and directions ought to be, in order to achieve a desired acceptance. The purpose of using an adaptive MCMC over the “vanilla” MCMC version shown in the previous chapters on MCMC because 1) you will get the opportunity to learn both versions and 2) speed is a priority for this lab exercise and using adaptation can lessen the number of iterations needed to show evidence of convergence.

In R, the package `adaptMCMC` includes the `MCMC` function, which implements adaptive MCMC. For this function, you must specify several arguments. First, you will specify your desired acceptance rate and set `adapt` to `TRUE` so adaptive sampling is used. You learned from the first chapter on MCMC that “the optimal acceptance rate varies depending on the number of dimensions in the parameter space.” You also learned that for a single parameter the optimal acceptance rate is about 44%, but as the number of parameters increases, the acceptance rate stabilizes around 23.4%. For the sea-level rise model, there are in total five parameters; three model parameters (α , T_{eq} , SL_0) and two statistical parameters (σ , and ρ). Hence, the desired acceptance rate is likely around 23.4%. Second, set the rate of adaptation `gamma.mcmc` to equal 0.5; this is because a lower γ leads to a faster adaptation. Next, you will set the step size based on step sizes used in the calibration of the Rahmstorf model in Ruckert et al. 2017. Even though the step sizes adapt, better initial estimates of the proposal covariance matrix yield faster convergence of the algorithm. In practice, you can play around with the “vanilla” MCMC function to observe how changing the step size for each parameter impacts the acceptance rate and use the results to inform the initial step sizes for the adaptive version. Lastly, you will set the MCMC analysis to run for 4×10^5 iterations and set the analysis to start adapting the step size after 1% of the iteration has run. Note that you do not want to start adapting too early because estimates of the parameter correlations will be weak if not enough proposals have been accepted yet. This analysis may take several minutes to complete.

```
# Set optimal acceptance rate as # parameters->infinity
# (Gelman et al, 1996; Roberts et al, 1997)
# Set number of iterations and rate of adaptation
# (gamma.mcmc, between 0.5 and 1, lower is faster adaptation)
accept.mcmc = 0.234
gamma.mcmc = 0.5
NI <- 4e5
step <- c(0.2, 0.02, 1, 0.1, 0.01)

# Run MCMC calibration. Note: Runs for several minutes.
mcmc.out <- MCMC(log.post, NI, p0, scale = step, adapt = TRUE, acc.rate = accept.mcmc,
                gamma = gamma.mcmc, list = TRUE, n.start = round(0.01*NI))
mcmc.chains <- mcmc.out$samples
```

Determine the burn-in period

In the previous chapter, you applied a 1% burn-in period, but how do you decide the appropriate length of the burn-in period? Besides visual inspection, you could use the Heidelberger and Welch diagnostic or the Gelman and Rubin diagnostic. The Gelman and Rubin diagnostic technique is useful because it can also be used to check for evidence of convergence as was described in the previous chapter on MCMC. First, you will run the MCMC analysis two more times and check for convergence at different points along the chains. Once the chains show evidence for convergence, then you can determine the burn-in. For this, look at how many iterations are required for the Gelman and Rubin potential scale reduction factor to get and stay below 1.1 for all parameters. It is only after the chains are converged that you should use the parameter values as posterior draws, for analysis. Therefore, the iterations prior to the point at which the scale reduction factor falls below 1.1 is what is thrown away as the burn-in period.

```

## Gelman and Rubin's convergence diagnostic:
set.seed(1708)
p0 <- c(1.9, -0.9, -145, 4, 0.7) # Arbitrary choice.
mcmc.out2 <- MCMC(log.post, NI, p0, scale = step, adapt = TRUE, acc.rate = accept.mcmc,
                 gamma = gamma.mcmc, list = TRUE, n.start = round(0.01*NI))
mcmc.chains2 <- mcmc.out2$samples

set.seed(1234)
p0 <- c(2.9, 0, -160, 5, 0.8) # Arbitrary choice.
mcmc.out3 <- MCMC(log.post, NI, p0, scale = step, adapt = TRUE, acc.rate = accept.mcmc,
                 gamma = gamma.mcmc, list = TRUE, n.start = round(0.01*NI))
mcmc.chains3 <- mcmc.out3$samples

# Turn the chains into an mcmc object.
# This is a requirement for the gelman.diag and gelman.plot command.
mcmc1 <- as.mcmc(mcmc.chains)
mcmc2 <- as.mcmc(mcmc.chains2)
mcmc3 <- as.mcmc(mcmc.chains3)

# Test for convergence.
set.seed(111) # revert back to original seed
mcmc_chain_list <- mcmc.list(list(mcmc1, mcmc2, mcmc3))
gelman.diag(mcmc_chain_list)
gelman.plot(mcmc_chain_list)

# Create a vector to test the statistic at several spots throughout the chain.
# Test from 5000 to 4e5 in increments of 5000.
niter.test <- seq(from = 5000, to = NI, by = 5000)
gr.test <- rep(NA, length(niter.test))

# Once the statistic is close to 1 (adopting the standard of < 1.1), the between-chain
# variability is indistinguishable from the within-chain variability, and hence converged.
for (i in 1:length(niter.test)){
  mcmc1 = as.mcmc(mcmc.chains[1:niter.test[i],])
  mcmc2 = as.mcmc(mcmc.chains2[1:niter.test[i],])
  mcmc3 = as.mcmc(mcmc.chains3[1:niter.test[i],])
  mcmc_chain_list = mcmc.list(list(mcmc1, mcmc2, mcmc3))
  gr.test[i] = gelman.diag(mcmc_chain_list)[2]
}

# Plot Gelman and Rubin statistics as a function of iterations. The iterations prior
# to the point in which convergence happens is what we set as the burn-in.
plot(niter.test, gr.test, pch=20)
abline(h = 1.1, col = "red")
gr.test < 1.1; first.stat = max(which((gr.test < 1.1) == FALSE))+1; print(first.stat)

# Set the burn-in to the 1st statistic that remains under 1.1 and remove it.
burnin <- seq(1, niter.test[first.stat], 1)
slr.burnin.chain <- mcmc.chains[-burnin, ]

```

Hindcasting and projecting sea-level rise

The parameters from the chains (burn-in removed) can be plugged into the model to generate multiple simulations of the data, called an “ensemble”. The residuals can be added to the simulations of the data to generate hindcasts and projections for further analysis. In the exercise below, you will generate hindcasts and projections with this process.

If you insert the following lines at the bottom of your R script, then you will make sea-level rise projections as a matrix (`SLR.projections`), using the sea-level model and multiple `for` loops. In `SLR.projections`, each row represents a state of the world while the columns represent time. Because you are using all 4×10^5 iterations, this code block will take several minutes to fully generate the projections. For an additional exercise, rather than using multiple `for` loops try to improve on the code by using the `apply` or `sapply` function in R, or vectorize the internal `for` loop in the `SLR.residuals` calculation. Does using any of these other options increase the speed of the code?

```
# Extract parameter vectors from the chain to enhance code readability.
alpha.chain <- slr.burnin.chain[,1]
Teq.chain <- slr.burnin.chain[,2]
SL_0.chain <- slr.burnin.chain[,3]
sigma.chain <- slr.burnin.chain[,4]
rho.chain <- slr.burnin.chain[,5]

# Set up empty matrices for sea-level output.
# Loop over the sea level model to generate a distribution of sea level simulations.
NI_length <- length(slr.burnin.chain[,1])
SLR.model.sim <- mat.or.vec(NI_length, length(project.years))
for(n in 1:NI_length) {
  SLR.model.sim[n, ] = gmsl_r07(a = alpha.chain[n], Teq = Teq.chain[n],
                                SL0 = SL_0.chain[n], tstep, Tg = rcp85)
}

# Estimate the residuals with the AR(1) coefficient and sigma.
SLR.residuals <- mat.or.vec(NI_length, length(project.years))  #(nr,nc)
for(n in 1:NI_length) {
  for(i in 2:length(project.years)) {
    SLR.residuals[n,i] <- rho.chain[n]*SLR.residuals[n,i-1] +
      rnorm(1, mean = 0, sd = sigma.chain[n])
  }
}

# Estimate the hindcasts and projections: add the residuals onto the model simulations.
SLR.projections <- mat.or.vec(NI_length, length(project.years))  #(nr,nc)
for(i in 1:NI_length) {
  SLR.projections[i,] <- SLR.model.sim[i,] + SLR.residuals[i,]
}
```

Exercise

Part 1. Using the diagnostics specified in the previous chapter, check the analysis for evidence of convergence. In your lab report, provide evidence to support or reject whether the chains are converged. If the chains are not converged, then explain how you would go about getting the chains to converge. The goal of this chapter is to understand how to apply a Markov chain Monte Carlo analysis to “real” data and use a “real” model, so getting the chains to converge—while extremely important in a physical application—is not a top priority in this chapter due to time constraints.




Figure 7: The sea level anomaly data from Church and White (2011; left) with its associated 90% credible interval, and the 90% credible interval projections of sea level anomalies from 2000 to 2100 in mm (right).

Part 2. Modify your script from Part 1 to:

- calculate and plot the 90% credible interval of the sea-level projections (to the year 2100) using the RCP8.5 temperature scenario (hint: consider using `quantile()` and `polygon()`),
- and to make density plots of the parameters α , T_{eq} , SL_0 , and ρ .

It is highly recommended that you do not plot all 4×10^5 projections as individual lines at once, but rather plot credible intervals. Ultimately, you should end up with a plot similar to Figure 1. If not, check your script against the codes listed through the tutorial to make sure your script is working properly.

Part 3. Save a copy of your script from Part 2 using a different file name. In this script, modify the MCMC analysis to assume the errors are homoskedastic (constant through time) by changing `y.meas.err` to a vector where the values equal 0, `y.meas.err <- rep(0, length(slr))`. Setting the measurement errors to 0 combines the model error and observational error into one term.

Your scripts should make the following plots:

1. a two-panel plot displaying the autocorrelation function and the time dependent measurement error of the sea-level data.
2. a plot comparing the 90% credible interval of projections made from considering heteroskedasticity (part 2) versus assuming homoskedasticity (part 3).
3. a four-panel plot comparing the densities of parameters α , T_0 , H_0 , and ρ made from considering heteroskedasticity (part 2) versus assuming homoskedasticity (part 3).

Questions

1. Compare the sea-level projections from *part 2* and *part 3*. How do the results differ? How much larger or smaller is the 90% credible interval in *part 1* from *part 2*? How might the assumptions made about the data impact “real-life” decisions regarding flood protection?
2. Compare the parameter distributions from *part 2* and *part 3*. How do the results differ? Do both methods produce the same mode and the same width in distribution?
3. Refer to the paper Ruckert et al. (2017) (link) and the results from *part 2* and *part 3*. Does it matter how you account for the error structure of data? Explain your reasoning.

Appendix

Comparing un-whitened raw residuals to the whitened residuals In the exercise above, you whiten the residuals from the sea-level data in order to satisfy the MCMC assumptions. This process changes the residuals to be normally distributed. Using the initial residuals and *rho* calculate the whitened residuals and compare them to the un-whitened raw residuals and to a normal distribution. For this case, does whitening the residuals make a difference?

```
# Whiten the residuals
n <- length(res)
w_res <- res[2:n] - start_rho*res[1:(n-1)]

par(mfrow = c(1,1))
plot(density(w_res), col="black", xlab="Residuals", yaxt="n", main="")
lines(density(res), col="red")
lines(density(rnorm(1e5, sd=5)), col="gray", lty=2)
abline(v=0, col="gray", lty=2)
legend("topleft", c("Whitened", "Un-whitened", "Normal distribution\nwith mean 0"),
      lty=c(1,1,2), col=c("black", "red", "gray"), bty="n")
```

Using different prior distributions Modify the prior parameter distributions (priors do not have to be uniform distributions). Explain your reasoning behind the modifications and rerun the MCMC calibration. How do the results differ?

Thinning the chain *Thinning* the chain is the process of taking a subsample of the chain instead of all of the chain. For example, you can take every “n’t” observation from the chain. Note that thinning the chain is optional and often not necessary. Here, we give an example of how it is done in practice.

Thinning can be useful for two reasons: (1) to reduce storage burden and (2) to increase computational speed. A good way to inform the choice of thinning is to look at the autocorrelation function of the Markov chains. We adopt the standard that the chains should be no smaller in length than the lag at which the autocorrelation function is below 0.05 or (5%), so the thinned parameters are not substantially autocorrelated (Link and Eaton, 2012). The plots show the parameters are relatively uncorrelated by roughly the 5,000th lag (this number may vary because the chains are not converged). This means that to properly thin the chain only every 5,000th iteration (or more) should be saved and that the samples are in general highly correlated with one another.

```
# The diagonal ACF plots are the plots of interest
par(mfrow = c(3,2))
for(i in 1:5){
  acf(slr.burnin.chain[,i], lag.max = 20000,
      main = paste('Parameter = ', parnames[i], sep = ''))
  abline(h = 0.05, col = "red")
}
```

When you thin the chains, visually check to make sure the full chain and the subset have almost identical posterior parameter distributions, so no important information is lost. The plot created below shows that the distributions are similar in shape, however, they do not visually match in shape.

```
# Thin the chain to a subset by taking every 5000th iteration (this number
# may vary if the chains are not converged). Another option is to
# use the sample() command.
subset_length <- round(nrow(slr.burnin.chain)/5000, 0)
sub_chain = slr.burnin.chain[seq(1, nrow(slr.burnin.chain), 5000), ]
# sub_chain2 <- slr.burnin.chain[sample(nrow(slr.burnin.chain), size = subset_length,
#                                     replace = FALSE), ]
```

```
# Check for similarities between full chain and the subset.
par(mfrow = c(3,2))
for(i in 1:5){
  plot(density(slr.burnin.chain[ ,i]), type="l",
       xlab = paste('Parameter =', ' ', parnames[i], sep=' '), ylab="PDF", main="")
  lines(density(sub_chain[ ,i]), col="red")
}
```

In this chapter, the output was not thinned. Why was thinning not necessary as well as inappropriate for this example?

####Applying MCMC to other data and models In previous exercises, you used the sea-level data from Jevrejeva et al (2008). Apply MCMC analysis accounting for autocorrelation and heteroskedastic residuals on the Jevrejeva et al (2008) data. For this, use the second order polynomial from earlier chapters as the model to fit to the data. You can extract the third vector from `sl.data` as your measurement error, `err.sl <- sl.data[,3]` and do not forget to set `y.meas.err <- err.sl`. Examine your sea-level rise estimate in 2100. How do your estimates differ from Exercise #7? How do your model parameter distributions differ from Exercise #7?

####Bootstrapping the sea-level model Apply the Bootstrap method to the Rahmstorf (2007) model with the data used in this chapter. How would one account for autocorrelation? Can one account for heteroskedasticity using the Bootstrap method?

References

- Bayes, T. 1764. An essay toward solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society on London* 53, 370-418. Available online at <http://rstl.royalsocietypublishing.org/content/53/370.full.pdf+html>
- Church, J. A. and White, N. J. 2011. Sea-Level Rise from the Late 19th to the Early 21st Century. *Surv Geophys*, 32, 585-602. Available online at <https://link.springer.com/article/10.1007/s10712-011-9119-1>.
- Gelman, A. and Rubin, D. B. 1992. Inference from iterative simulation using multiple sequences. *Statistical Science*, 7, 457-511. Available online at <https://projecteuclid.org/euclid.ss/1177011136>
- Gilks, W. R. 1997. *Markov chain Monte Carlo in practice*. Chapman & Hall/CRC Press, London, UK
- Hastings, W. K. 1970. Monte Carlo sampling methods using Markov chains and their applications. *Biometrika* 57(1):97-109. doi:10.1093/biomet/57.1.97. Available online at https://www.jstor.org/stable/2334940?seq=1#page_scan_tab_contents
- Jevrejeva, S., Moore, J. C., Grinsted, A., and Woodworth, P. L., 2008. Recent global sea level acceleration started over 200 years ago? *Geophysical Research Letters* 35, L08715. Available online at <http://onlinelibrary.wiley.com/doi/10.1029/2008GL033611/full>
- Link, W. A. and Eaton, M. J. 2012. On thinning of chains in MCMC. *Methods in Ecology and Evolution* 3, 112-115. Available online at <http://onlinelibrary.wiley.com/doi/10.1111/j.2041-210X.2011.00131.x/abstract>
- Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., and Teller, E. Equation of state calculations by fast computing machines. *J. Chem. Phys.* 21(6), 1087-1092. Available online at <http://aip.scitation.org/doi/abs/10.1063/1.1699114>
- Rahmstorf, S. 2007. A Semi-empirical approach to projecting future sea-level rise. *Science* 315, 368-370. Available online at <http://science.sciencemag.org/content/315/5810/368>
- Roberts, G. O., Gelman, A., and Gilks W. R. 1997. Weak convergence and optimal scaling of random walk Metropolis algorithms. *Ann. Appl. Prob.* 7, 110-120. Available online at http://projecteuclid.org/download/pdf_1/euclid.aoap/1034625254

- Rosenthal, J. S. 2010. "Optimal Proposal Distributions and Adaptive MCMC." Handbook of Markov chain Monte Carlo. Eds., Brooks, S., Gelman, A., Jones, G. L., and Meng, X.-L. Chapman & Hall/CRC Press. Available online at <https://pdfs.semanticscholar.org/3576/ee874e983908f9214318abb8ca425316c9ed.pdf>
- Ruckert, K. L., Guan, Y., Bakker, A. M. R., Forest, C. E., and Keller, K. The effects of non-stationary observation errors on semi-empirical sea-level projections. *Climatic Change* 140(3), 349-360. Available online at <http://dx.doi.org/10.1007/s10584-016-1858-z>
- Smith, T. M., Reynolds, R. W., Peterson, T. C., and Lawrimore, J. 2008. Improvements to NOAA's historical merged land-ocean surface temperature analysis (1880-2006). *J. Clim.* 21, 2283–2296. Available online at <http://journals.ametsoc.org/doi/abs/10.1175/2007JCLI2100.1>
- Vihola, M. 2011. Robust adaptive Metropolis algorithm with coerced acceptance rate. *Statistics and Computing*. Available online at <http://www.springerlink.com/content/672270222w79h431/>
- Zellner, A. and Tiao, G. C. 1964. Bayesian analysis of the regression model with autocorrelated errors. *J. Am. Stat. Assoc.* 59, 763–778. available online at <http://www.sciencedirect.com/science/article/pii/S0167947301000846>

Chapter #5: Global sensitivity analyses (Tony E. Wong, Vivek Srikrishnan and Klaus Keller)

Chapter #6: Decision-making under uncertainty (Vivek Srikrishnan and Klaus Keller)

Chapter #7: Climate policy and the DICE model (Patrick J. Applegate, Gregory G. Garner, Richard B. Alley, and Klaus Keller)

Learning objectives

After completing this exercise, you should be able to

- describe what the social cost of carbon (SCC) is
- explain in broad terms what the DICE model is and how it works
- perform simple calculations with the DICE model
- describe how uncertainties in the climate sensitivity affect the present-day social cost of carbon

Introduction

As noted in the Introduction to this e-textbook, fossil fuel use creates benefits for people, but also imposes costs in the form of climate change. Burning fossil fuels releases energy and waste products including carbon dioxide and water. The energy leads to economic productivity. However, the carbon dioxide causes long-lasting temperature increases, which lead to increased risks for people in the future. Most of the impacts associated with climate change will probably be harmful for people.

Because fossil fuel users pay less than the full cost associated with their actions, emissions of carbon dioxide are larger than the economically-optimal amount. An individual fossil fuel user pays an amount of money per unit of energy that represents the cost of extraction, processing (for example, gasoline is a highly processed derivative of oil), transportation, and profits, plus taxes and minus any subsidies. This cost does not reflect the increased risks to future people associated with climate change. Because people respond to price signals, they tend to consume more fossil fuel than they would if the price reflected the full cost to society of fossil fuel consumption (and were therefore higher).

The social cost of carbon (SCC) is the increase in climate-related damages caused by emitting an additional ton of CO₂ to the atmosphere (Nordhaus, 2013). In other words, if we were to emit one ton of CO₂ today, the SCC would be the sum of the negative impacts caused by that additional ton of CO₂, from now into the distant future. The future damages from carbon dioxide emissions are typically discounted in calculating the SCC; see Exercise 4 for a brief discussion of discounting. In 2010, the US Environmental Protection Agency estimated a value for the SCC of about \$21/t CO₂ (Interagency Working Group on Social Cost of Carbon, 2010, 2013; Johnson and Hope, 2012).

The Dynamic Integrated model of Climate and the Economy (DICE; Nordhaus, 2013; Nordhaus and Sztorc, 2013) includes the feedbacks between the climate system and the economy that are needed for estimating the SCC. As discussed previously, greenhouse gas emissions cause temperature increases, which cause climate-related damages to the economy. These damages can then motivate the development of policies and technologies that reduce emissions and possibly remove greenhouse gases from the atmosphere. A model like DICE that “closes the loop” between climate change and further emissions is necessary for accurate estimation of future damages and therefore the SCC. DICE is one of three models used by the EPA to estimate the social cost of carbon.

DICE is commonly used to identify optimal climate policies, sometimes subject to different constraints. When run in optimization mode, DICE maximizes a “utility function” that includes discounting as well as the declining marginal benefit of additional consumption. That is, rich societies benefit less from an additional dollar of income than poor ones.

However, before DICE can be optimized, some assumptions have to be made about the climate sensitivity, among other parameters. As we saw in the Introduction, the climate sensitivity represents the amount by which global mean air temperatures would increase if carbon dioxide concentrations in the atmosphere were to double. Although many studies have estimated the climate sensitivity, this parameter remains deeply uncertain.

In this exercise, we use DICE to investigate how different climate and economic variables might change in the future under an optimized climate policy, plus how different values of the climate sensitivity affect the present value of the SCC.

An important note

Carrying out this exercise will not prepare you to estimate the social cost of carbon for publication in scientific journals or for use in policy applications. The presentation of the material in this exercise has been simplified for pedagogical reasons, and this material should only be applied to “real” problems after additional training. The authors and editors of this e-textbook specifically disclaim any liability for damages associated with the use of the material presented in this exercise.

Tutorial

The version of the DICE model that we use here was translated to R by Greg Garner from an original version in the GAMS language provided by Bill Nordhaus. The `dice.R` file contains the following useful functions (among others):

- `dice.new()`: creating a new instance of the DICE model with a prechosen set of parameters
- `dice.modify()`: changing the value of model parameters in DICE
- `dice.run()`: running the DICE model into the future using a given set of parameters
- `dice.solve()`: identifying an optimal trajectory using the utility function sketched above

Open `lab8_sample.R` in RStudio and examine the code that it contains. The following lines of code explain who wrote the code, describe what the code does, `source()`s the file containing the DICE model, and loads the DICE model into memory using the `dice.new()` function.

```
# lab8_done.R  
# Patrick Applegate, patrick.applegate@psu.edu; Greg Garner, ggg121@psu.edu  
#  
# Optimizes the DICE model to produce a plausible climate-economic trajectory  
# and performs a Monte Carlo experiment to evaluate the effects of uncertainty  
# in the climate sensitivity on the present-day social cost of carbon.  
  
# Mise en place.  
rm(list = ls())  
graphics.off()  
  
# Load the DICE model.  
source("dice.R")  
  
# Create a new DICE object.  
my.dice <- dice.new()
```

Next, we optimize the DICE model using the utility function described in the Introduction and pull some time-dependent output out of the object that results.

```
# Solve the DICE object for the optimal control policy.  
# NOTE: THIS STEP MAY TAKE A COUPLE OF MINUTES!  
dice.solve(my.dice)  
  
# Set sensible names for the time-dependent variables to extract from DICE.  
names <- c("Time (yr)",  
          "Emissions (Gt CO2/yr)",  
          "Atmospheric [CO2] (ppm)",
```

```

    "Global mean T anomaly (C)",
    "Climate damages (1012 $)",
    "Social cost of carbon ($/t CO2)")

# Make a place to store time-dependent output from optimized DICE.
opt.output <- matrix(data = NA, nrow = length(my.dice$year), ncol = 6)
colnames(opt.output) <- names

# Put the time-dependent output from optimized DICE into opt.output.
opt.output[, 1] <- my.dice$year
opt.output[, 2] <- my.dice$e
opt.output[, 3] <- my.dice$mat
opt.output[, 4] <- my.dice$statm
opt.output[, 5] <- my.dice$damages
opt.output[, 6] <- my.dice$scc

```

The script also makes a plot of the time-dependent output.

```

# Plot the time-dependent output from the optimized DICE object.
pdf("lab8_plot1.pdf", width = 5, height = 8.5)
par(mfrow = c(5, 1))
plot(opt.output[, 1], opt.output[, 2], type = "l", bty = "n", xlab = names[1],
      ylab = names[2])
plot(opt.output[, 1], opt.output[, 3], type = "l", bty = "n", xlab = names[1],
      ylab = names[3])
plot(opt.output[, 1], opt.output[, 4], type = "l", bty = "n", xlab = names[1],
      ylab = names[4])
plot(opt.output[, 1], opt.output[, 5], type = "l", bty = "n", xlab = names[1],
      ylab = names[5])
plot(opt.output[, 1], opt.output[, 6], type = "l", bty = "n", xlab = names[1],
      ylab = names[6])
dev.off()

```

Running DICE in optimization mode and examining its time-varying output

If you haven't already done so, download the .zip file containing the scripts associated with this book from www.scrimhub.org/raes. Put the file `lab8_sample.R` in an empty directory. Execute `lab8_sample.R` using the `source()` command or button (remember to set your working directory first), and look at the .pdf file that results. It should look like Figure 14.

The script also prints out the assumed climate sensitivity value and the present-day social cost of carbon (`opt.scc <- opt.output[1, 6]`) obtained using the optimization.

```

# Also extract the preset climate sensitivity and social cost of carbon
# from the DICE object.
opt.t2xC02 <- my.dice$t2xC02
opt.scc <- opt.output[1, 6]

# Print some key quantities.
print(sprintf("The assumed climate sensitivity for optimization is %2.2f C/doubling",
              opt.t2xC02))
print(sprintf("The optimized social cost of carbon in %d is $%4.2f/t CO2",
              opt.output[1, 1], opt.scc))

```

Several interesting observations come out of Figure 14:




Figure 8: Time-dependent output from the optimized DICE model using the built-in parameter set. See text for discussion.

1. Each model curve bends sharply upward in the last few decades.
2. Ignoring the last few decades of model results, the peak of each curve is lagged in time relative to the peak of the preceding curve; so, peak temperatures occur after peak carbon dioxide concentrations, which occur after peak emissions.
3. The social cost of carbon is not constant; it starts out relatively small and grows over time until about 2200, when it begins to decline again.

We can ignore the last few decades of each time series shown in Figure 14. The DICE model assumes, in effect, that the world ends shortly after 2300 and that this apocalypse is known in advance, so that the people of the world start burning fossil fuels again.

The lags between the peaks in each curve happen for the following reasons.

- The atmospheric CO₂ concentration continues to increase after maximum net CO₂ emissions because the rate at which CO₂ is removed from the atmosphere is still smaller than the rate of release. Once net CO₂ emissions balance with the rate of CO₂ removal from the atmosphere, the atmospheric CO₂ peaks and subsequently declines as the net CO₂ emissions continue to drop.
- The atmospheric temperature peaks later than the atmospheric CO₂ due to the way energy (in the form of heat) is moved between the atmosphere and the ocean. Although the atmosphere responds quickly to additional forcing, the oceans are slow to remove the heat from the atmosphere. The peak in atmospheric temperature occurs when the ocean heat uptake balances the radiative forcing produced by the level of CO₂ in the atmosphere.
- The climate damages are a function of atmospheric temperature and gross world product. Once atmospheric temperatures begin to decrease, the damages as a fraction of the gross world product decreases; however, the gross world product is constantly increasing over time, and this continued increase results in the delayed peak in climate damages.
- As discussed earlier, the social cost of carbon represents the sum of future climate damages from the emission of a single ton of CO₂. Because the model assumes that the world ends shortly after 2300, the future damages due to a unit increase in CO₂ emissions decline as the model simulation approaches the end of the simulation period. As a result, the social cost of carbon first increases, then decreases.

Accounting for uncertainty in the climate sensitivity

DICE assumes a climate sensitivity value of 2.9 C per doubling of CO₂ concentration in the atmosphere. Although this value is reasonable, it doesn't capture our uncertainty in the actual value of the climate sensitivity. If we were to change this parameter within reasonable limits, we would get different climate and economic trajectories (Fig. 14), and a different present-day value for the social cost of carbon.

We could account for the effects of uncertainty in the climate sensitivity on present-day social cost of carbon values by performing a Monte Carlo experiment; however, we would first need a well-defined distribution for the climate sensitivity. We could then run DICE repeatedly with climate sensitivity values sampled from this distribution and examine the distribution of present-day social cost of carbon values that came out.

One *approximate* method involves matching a lognormal distribution to the probabilistic statements of the Intergovernmental Panel on Climate Change on climate sensitivity. The lognormal distribution is only defined for positive values, and is therefore appropriate for variables like climate sensitivity that can't be negative. The latest IPCC report says, "Equilibrium climate sensitivity is... extremely unlikely less than 1°C... and very unlikely greater than 6°C..." (Alexander et al., 2013). Mastrandrea et al. (2010, their Table 1) suggest that "extremely unlikely" corresponds to a likelihood of 0-5%, and "very unlikely" corresponds to a likelihood of 0-10%.

The following code block accomplishes this matching, assuming the largest likelihood values for "extremely unlikely" and "very unlikely" (5% and 10%, respectively). It also generates a vector of random values from the resulting distribution.

```
# Set some values.
xs <- c(1, 6)      # climate sensitivities corresponding to the probabilities
```

```

# in ps
ps <- c(0.05, (1- 0.1))
# (approximate) probabilities of the climate sensitivity
# being less than xs[1] or greater than xs[2], according
# to IPCC AR5 WG1
n.trials <- 300 # number of Monte Carlo trials

# Define a function for matching the lognormal distribution to two (or more)
# tie points.
lnorm.rmse <- function(dist.params, xs, ps) {
  # dist.params, vector of meanlog and sdlog values;
  # see help(dlnorm)
  # xs, vector of values of the distributed variable to match
  # ps, probabilities of the values in xs
  logmu <- dist.params[1]
  logsigma <- dist.params[2]
  trial.xs <- qlnorm(ps, logmu, logsigma)
  rmse <- sqrt(mean((xs- trial.xs)^ 2))
  return(rmse)
}

# Identify the parameters of the lognormal distribution that produce the best
# match to the IPCC's statements about climate sensitivity.
lnorm.optim <- optim(log(c(2.9, 1.5)), lnorm.rmse, gr = NULL, xs = xs, ps = ps,
  method = "L-BFGS-B", lower = c(0, 0), upper = c(Inf, Inf))

# Generate a vector of climate sensitivity values.
set.seed(1)
tx2C02s <- rlnorm(n.trials, meanlog = lnorm.optim$par[1],
  sdlog = lnorm.optim$par[2])

```

Exercise

Make a copy of `lab8_sample.R` by saving it under a different file name. Modify this copy so that it performs the Monte Carlo experiment described above using a `for` loop, with `n.trials <- 300` samples. For each value in `tx2C02s`, you'll need to

1. set the value of climate sensitivity in DICE using `dice.modify(my.dice, "t2xC02", tx2C02s[i])`
2. run the DICE model using `dice.run(my.dice)`
3. extract the 2010 value of the social cost of carbon from the DICE object and store it in the `i`th element of a vector `sccs`

Your modified script should make a plot with two panels. The top panel should show the distribution of climate sensitivity values based on fitting a lognormal distribution to the IPCC's probabilistic statement, with a vertical line to show DICE's default climate sensitivity value. The bottom panel should show the distribution of present-day social cost of carbon values that you obtained from your Monte Carlo experiment, with vertical lines showing the mean of these values and the social cost of carbon value from optimizing DICE with the default parameter values.

Questions

1. The mean social cost of carbon from your Monte Carlo experiment reflects an estimate of the social cost of carbon given our remaining uncertainty in climate sensitivity. Is this value higher or lower than

- the social cost of carbon value from optimizing DICE with the base parameters? How much higher or lower is it? Express your answer as a percentage.
2. How sensitive is your answer to question 1 to the likelihoods you assign to “extremely unlikely” and “very unlikely,” given the likelihood ranges specified for these terms by Mastrandrea et al (2010, their Table 1)? Does one bound have more of an effect on the mean social cost of carbon from your Monte Carlo experiments than the other?
 3. If we are uncertain about the actual value of climate sensitivity, should we spend more or less money on reducing carbon dioxide emissions now, compared to a case in which we are sure about the value of climate sensitivity? Justify your answer based on your responses to questions 1 and 2.

References

- Alexander, L. V., et al., 2013. Summary for policymakers. In Stocker, T. F., et al., eds., *Climate Change 2013: The Physical Science Basis. Contribution of Working Group I to the Fifth Assessment Report of the Intergovernmental Panel on Climate Change*, Cambridge University Press. Available online at <http://www.climatechange2013.org/report/full-report/>
- Interagency Working Group on Social Cost of Carbon, 2010. Technical Support Document: Social Cost of Carbon for Regulatory Impact Analysis Under Executive Order 12866. Available online at www.epa.gov/otaq/climate/regulations/scc-tsd.pdf
- Interagency Working Group on Social Cost of Carbon, 2013. Technical Support Document: Technical Update of the Social Cost of Carbon for Regulatory Impact Analysis Under Executive Order 12866. Available online at https://www.whitehouse.gov/sites/default/files/omb/inforeg/social_cost_of_carbon_for_ria_2013_update.pdf
- Johnson, L. T., and Hope, C., 2012. The social cost of carbon in U.S. regulatory impact analyses: an introduction and critique. *Journal of Environmental Studies and Science* 2, 205-221, doi: 10.1007/s13412-012-0087-7. Available online at <http://link.springer.com/article/10.1007/s13412-012-0087-7>
- Mastrandrea, M. D., et al., 2010. Guidance Note for Lead Authors of the IPCC Fifth Assessment Report on Consistent Treatment of Uncertainties. Available online at <https://www.ipcc.ch/pdf/supporting-material/uncertainty-guidance-note.pdf>
- Nordhaus, W., 2013. *The Climate Casino: Risk, Uncertainty, and Economics for a Warming World*. Yale University Press, 378 p.
- Nordhaus, W., and Sztorc, P., 2013. DICE 2013R: Introduction and User’s Manual (2nd ed). Available online at http://www.econ.yale.edu/~nordhaus/homepage/documents/DICE_Manual_103113r2.pdf

Chapter #8: Multi-objective decision analysis with a lake management problem (Caitlin Spence, Vivek Srikrishnan, and Giacomo Marangoni)

Chapter #9: Introduction to multi-objective evolutionary algorithms and multi-objective optimization (Vivek Srikrishnan, Caitlin Spence, Giacomo Marangoni, and Patrick Applegate)

**Chapter #10: Vulnerability analysis and robust decision-making
(Vivek Srikrishnan, Caitlin Spence, and Giacomo Marangoni)**

**Chapter #11: Principles of data visualization (Vivek Srikrishnan,
Kelsey L. Ruckert, and Klaus Keller)**

Chapter #12: Visualizing uncertainty using sea-level rise and storm surge (Vivek Srikrishnan, Kelsey L. Ruckert, and Klaus Keller)

GNU GENERAL PUBLIC LICENSE Version 3 (Free Software Foundation, Inc.)

Copyright (C) 2007 Free Software Foundation, Inc. <http://fsf.org/> Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The GNU General Public License is a free, copyleft license for software and other kinds of works.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users. We, the Free Software Foundation, use the GNU General Public License for most of our software; it applies also to any other work released this way by its authors. You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

To protect your rights, we need to prevent others from denying you these rights or asking you to surrender the rights. Therefore, you have certain responsibilities if you distribute copies of the software, or if you modify it: responsibilities to respect the freedom of others.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must pass on to the recipients the same freedoms that you received. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

Developers that use the GNU GPL protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License giving you legal permission to copy, distribute and/or modify it.

For the developers' and authors' protection, the GPL clearly explains that there is no warranty for this free software. For both users' and authors' sake, the GPL requires that modified versions be marked as changed, so that their problems will not be attributed erroneously to authors of previous versions.

Some devices are designed to deny users access to install or run modified versions of the software inside them, although the manufacturer can do so. This is fundamentally incompatible with the aim of protecting users' freedom to change the software. The systematic pattern of such abuse occurs in the area of products for individuals to use, which is precisely where it is most unacceptable. Therefore, we have designed this version of the GPL to prohibit the practice for those products. If such problems arise substantially in other domains, we stand ready to extend this provision to those domains in future versions of the GPL, as needed to protect the freedom of users.

Finally, every program is threatened constantly by software patents. States should not allow patents to restrict development and use of software on general-purpose computers, but in those that do, we wish to avoid the special danger that patents applied to a free program could make it effectively proprietary. To prevent this, the GPL assures that patents cannot be used to render the program non-free.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS

0. Definitions.

“This License” refers to version 3 of the GNU General Public License.

“Copyright” also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

“The Program” refers to any copyrightable work licensed under this License. Each licensee is addressed as “you”. “Licensees” and “recipients” may be individuals or organizations.

To “modify” a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a “modified version” of the earlier work or a work “based on” the earlier work.

A “covered work” means either the unmodified Program or a work based on the Program.

To “propagate” a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To “convey” a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays “Appropriate Legal Notices” to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The “source code” for a work means the preferred form of the work for making modifications to it. “Object code” means any non-source form of a work.

A “Standard Interface” means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The “System Libraries” of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A “Major Component”, in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The “Corresponding Source” for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work’s System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of

having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users' Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

a) The work must carry prominent notices stating that you modified it, and giving a relevant date.

b) The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to "keep intact all notices".

c) You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.

d) If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or

on a volume of a storage or distribution medium, is called an “aggregate” if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation’s users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

a) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.

b) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.

c) Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.

d) Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.

e) Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A “User Product” is either (1) a “consumer product”, which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation

into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, “normally used” refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

“Installation Information” for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source. The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

7. Additional Terms.

“Additional permissions” are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- a) Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- b) Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or
- c) Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or
- d) Limiting the use for publicity purposes of names of licensors or

authors of the material; or

e) Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or

f) Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered “further restrictions” within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance. However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An “entity transaction” is a transaction transferring control of an organization, or substantially all assets of

one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party's predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A "contributor" is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor's "contributor version".

A contributor's "essential patent claims" are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, "control" includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor's essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a "patent license" is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To "grant" such a patent license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. "Knowingly relying" means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient's use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is "discriminatory" if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others' Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Use with the GNU Affero General Public License.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU Affero General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the special requirements of the GNU Affero General Public License, section 13, concerning interaction through a network will apply to the combination as such.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU General Public License “or any later version” applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU General Public License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR

A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.